

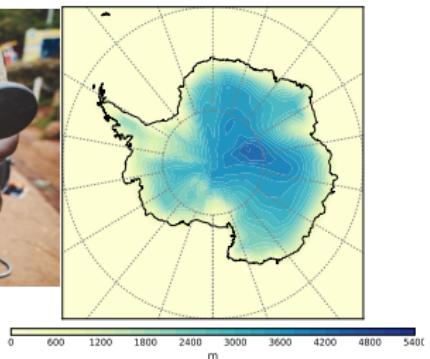
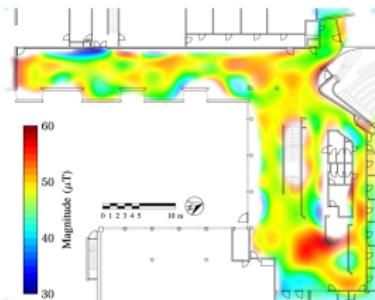
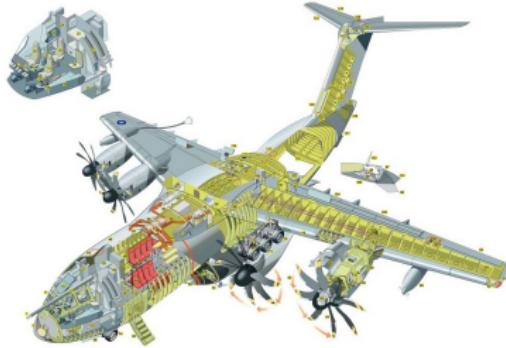
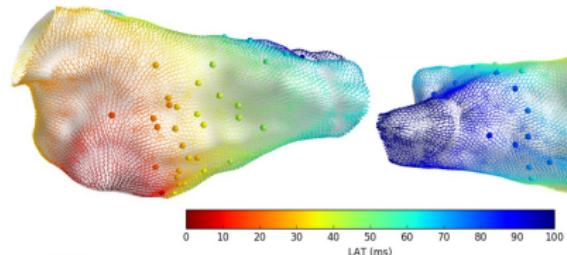
Gaussian Process Emulators for Cardiac Digital Twins

Richard Wilkinson, Chris Lanyon

School of Mathematical Sciences
University of Nottingham

Cardiac Digital Twin summer school
Oct 2025

Recent Gaussian process applications



Computer models

Some notation

Scientists often build complex models that simulate some phenomenon:

- x - inputs to the simulator
- $f(x)$ - outputs of the simulator

Computer models

Some notation

Scientists often build complex models that simulate some phenomenon:

- x - inputs to the simulator
- $f(x)$ - outputs of the simulator

We usually want to analyse the simulator in some way:

Computer models

Some notation

Scientists often build complex models that simulate some phenomenon:

- x - inputs to the simulator
- $f(x)$ - outputs of the simulator

We usually want to analyse the simulator in some way:

- Uncertainty analysis
 - ▶ Given a distribution for x , work out the distribution of $f(x)$

Computer models

Some notation

Scientists often build complex models that simulate some phenomenon:

- x - inputs to the simulator
- $f(x)$ - outputs of the simulator

We usually want to analyse the simulator in some way:

- Uncertainty analysis
 - ▶ Given a distribution for x , work out the distribution of $f(x)$
- Sensitivity analysis
 - ▶ Determine which input contributes most variance to the distribution of $f(x)$

Computer models

Some notation

Scientists often build complex models that simulate some phenomenon:

- x - inputs to the simulator
- $f(x)$ - outputs of the simulator

We usually want to analyse the simulator in some way:

- Uncertainty analysis
 - ▶ Given a distribution for x , work out the distribution of $f(x)$
- Sensitivity analysis
 - ▶ Determine which input contributes most variance to the distribution of $f(x)$
- Calibration
 - ▶ Given data $y = f(x^*) + \epsilon$, estimate the unknown parameter values x^*

Computer models

Some notation

Scientists often build complex models that simulate some phenomenon:

- x - inputs to the simulator
- $f(x)$ - outputs of the simulator

We usually want to analyse the simulator in some way:

- Uncertainty analysis
 - ▶ Given a distribution for x , work out the distribution of $f(x)$
- Sensitivity analysis
 - ▶ Determine which input contributes most variance to the distribution of $f(x)$
- Calibration
 - ▶ Given data $y = f(x^*) + \epsilon$, estimate the unknown parameter values x^*
- Calibrated prediction

Code uncertainty

Sacks, Welch, Mitchell, Wynn 1984

For complex simulators, run times might be long, making analysis challenging:

Code uncertainty

Sacks, Welch, Mitchell, Wynn 1984

For complex simulators, run times might be long, making analysis challenging:

- All inference must be done using a finite ensemble of model runs

$$\mathcal{D}_{sim} = \{(x_i, f(x_i))\}_{i=1,\dots,N}$$

- If x is not in the ensemble, then we are uncertain about the value of $f(x)$.

Code uncertainty

Sacks, Welch, Mitchell, Wynn 1984

For complex simulators, run times might be long, making analysis challenging:

- All inference must be done using a finite ensemble of model runs

$$\mathcal{D}_{sim} = \{(x_i, f(x_i))\}_{i=1,\dots,N}$$

- If x is not in the ensemble, then we are uncertain about the value of $f(x)$.

Idea: If the simulator is expensive, build a cheap model (*surrogate or emulator*) of it and use this in any analysis.

'a model of the model'

Code uncertainty

Sacks, Welch, Mitchell, Wynn 1984

For complex simulators, run times might be long, making analysis challenging:

- All inference must be done using a finite ensemble of model runs

$$\mathcal{D}_{sim} = \{(x_i, f(x_i))\}_{i=1,\dots,N}$$

- If x is not in the ensemble, then we are uncertain about the value of $f(x)$.

Idea: If the simulator is expensive, build a cheap model (*surrogate or emulator*) of it and use this in any analysis.

'a model of the model'

Gaussian processes (GPs) are the most widely used choice of emulator

Code uncertainty

Sacks, Welch, Mitchell, Wynn 1984

For complex simulators, run times might be long, making analysis challenging:

- All inference must be done using a finite ensemble of model runs

$$\mathcal{D}_{sim} = \{(x_i, f(x_i))\}_{i=1,\dots,N}$$

- If x is not in the ensemble, then we are uncertain about the value of $f(x)$.

Idea: If the simulator is expensive, build a cheap model (*surrogate or emulator*) of it and use this in any analysis.

'a model of the model'

Gaussian processes (GPs) are the most widely used choice of emulator

People also use polynomial chaos, neural networks, other ML/stats regression tools (eg linear regression, random forests), ...

Introduction

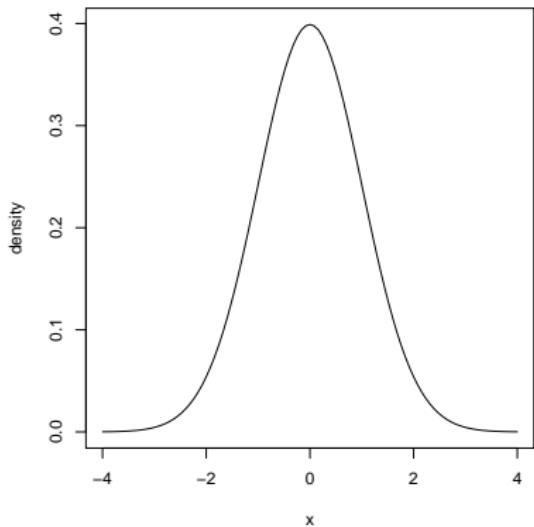
- Gaussian process basics
- Difficulties of working with GPs
- Example

You can download a copy of these slides from

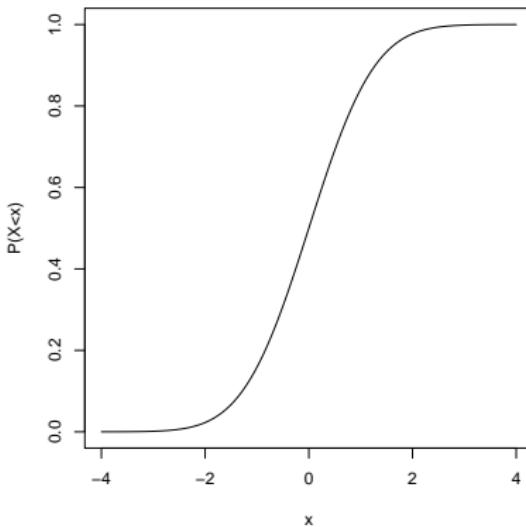
<https://github.com/cwlanyon/CardiacDTWorkshop>

Univariate Gaussian distributions

PDF of a $N(0,1)$ random variable



CDF of a $N(0,1)$ random variable



$$Y \sim N(\mu, \sigma^2)$$

PDF:

$$f_Y(y) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y-\mu)^2}{2\sigma^2}\right)$$

Univariate Gaussians

The normal/Gaussian distribution occurs naturally and is convenient mathematically

¹max. ent. principle: the distribution with the largest entropy should be used as a least-informative default

Univariate Gaussians

The normal/Gaussian distribution occurs naturally and is convenient mathematically

- Family of normal distributions is closed under linear operations.

¹max. ent. principle: the distribution with the largest entropy should be used as a least-informative default

Univariate Gaussians

The normal/Gaussian distribution occurs naturally and is convenient mathematically

- Family of normal distributions is closed under linear operations.
- Central limit theorem

¹max. ent. principle: the distribution with the largest entropy should be used as a least-informative default

Univariate Gaussians

The normal/Gaussian distribution occurs naturally and is convenient mathematically

- Family of normal distributions is closed under linear operations.
- Central limit theorem
- Maximum entropy/surprisal: $N(\mu, \sigma^2)$ has maximum entropy¹ of any distribution with mean μ and variance σ^2

¹max. ent. principle: the distribution with the largest entropy should be used as a least-informative default

Univariate Gaussians

The normal/Gaussian distribution occurs naturally and is convenient mathematically

- Family of normal distributions is closed under linear operations.
- Central limit theorem
- Maximum entropy/surprisal: $N(\mu, \sigma^2)$ has maximum entropy¹ of any distribution with mean μ and variance σ^2
- If Y and Z are jointly normally distributed and are uncorrelated, then they are independent

¹max. ent. principle: the distribution with the largest entropy should be used as a least-informative default

Univariate Gaussians

The normal/Gaussian distribution occurs naturally and is convenient mathematically

- Family of normal distributions is closed under linear operations.
- Central limit theorem
- Maximum entropy/surprisal: $N(\mu, \sigma^2)$ has maximum entropy¹ of any distribution with mean μ and variance σ^2
- If Y and Z are jointly normally distributed and are uncorrelated, then they are independent
- Square-loss functions lead to procedures that have a Gaussian probabilistic interpretation
eg minimizing $\sum(y_i - f_\beta(x_i))^2$

¹max. ent. principle: the distribution with the largest entropy should be used as a least-informative default

Multivariate Gaussian distributions

'Multivariate' = two or more random variables

Multivariate Gaussian distributions

'Multivariate' = two or more random variables

Suppose $Y \in \mathbb{R}^d$ has a multivariate Gaussian distribution with

- **mean vector** $\mu \in \mathbb{R}^d$
- **covariance matrix** $\Sigma \in \mathbb{R}^{d \times d}$.

Write

$$Y \sim N_d(\mu, \Sigma)$$

Multivariate Gaussian distributions

'Multivariate' = two or more random variables

Suppose $Y \in \mathbb{R}^d$ has a multivariate Gaussian distribution with

- **mean vector** $\mu \in \mathbb{R}^d$
- **covariance matrix** $\Sigma \in \mathbb{R}^{d \times d}$.

Write

$$Y \sim N_d(\mu, \Sigma)$$

Bivariate Gaussian: d=2

$$Y = \begin{pmatrix} Y_1 \\ Y_2 \end{pmatrix} \quad \mu = \begin{pmatrix} \mu_1 \\ \mu_2 \end{pmatrix} \quad \Sigma = \begin{pmatrix} \sigma_1^2 & \rho_{12}\sigma_1\sigma_2 \\ \rho_{21}\sigma_1\sigma_2 & \sigma_2^2 \end{pmatrix}$$

Multivariate Gaussian distributions

'Multivariate' = two or more random variables

Suppose $Y \in \mathbb{R}^d$ has a multivariate Gaussian distribution with

- **mean vector** $\mu \in \mathbb{R}^d$
- **covariance matrix** $\Sigma \in \mathbb{R}^{d \times d}$.

Write

$$Y \sim N_d(\mu, \Sigma)$$

Bivariate Gaussian: d=2

$$Y = \begin{pmatrix} Y_1 \\ Y_2 \end{pmatrix} \quad \mu = \begin{pmatrix} \mu_1 \\ \mu_2 \end{pmatrix} \quad \Sigma = \begin{pmatrix} \sigma_1^2 & \rho_{12}\sigma_1\sigma_2 \\ \rho_{21}\sigma_1\sigma_2 & \sigma_2^2 \end{pmatrix}$$

$$\text{Var}(Y_i) = \sigma_i^2 \quad \text{Cov}(Y_1, Y_2) = \rho_{12}\sigma_1\sigma_2 \quad \text{Cor}(Y_1, Y_2) = \rho_{12}$$

Multivariate Gaussian distributions

'Multivariate' = two or more random variables

Suppose $Y \in \mathbb{R}^d$ has a multivariate Gaussian distribution with

- **mean vector** $\mu \in \mathbb{R}^d$
- **covariance matrix** $\Sigma \in \mathbb{R}^{d \times d}$.

Write

$$Y \sim N_d(\mu, \Sigma)$$

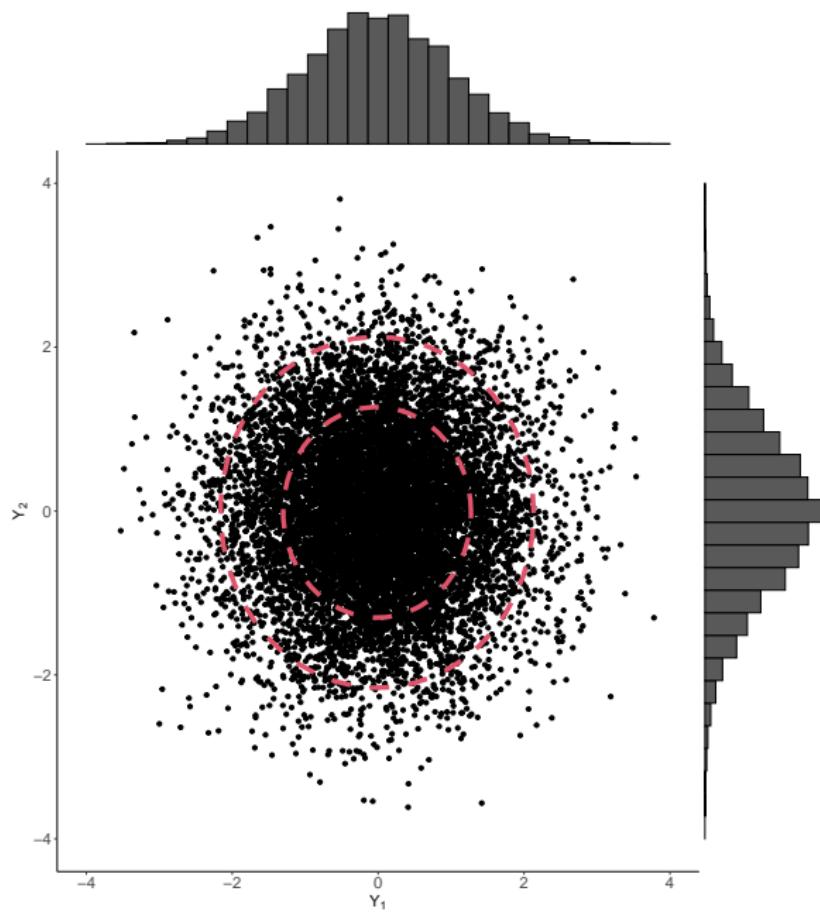
Bivariate Gaussian: $d=2$

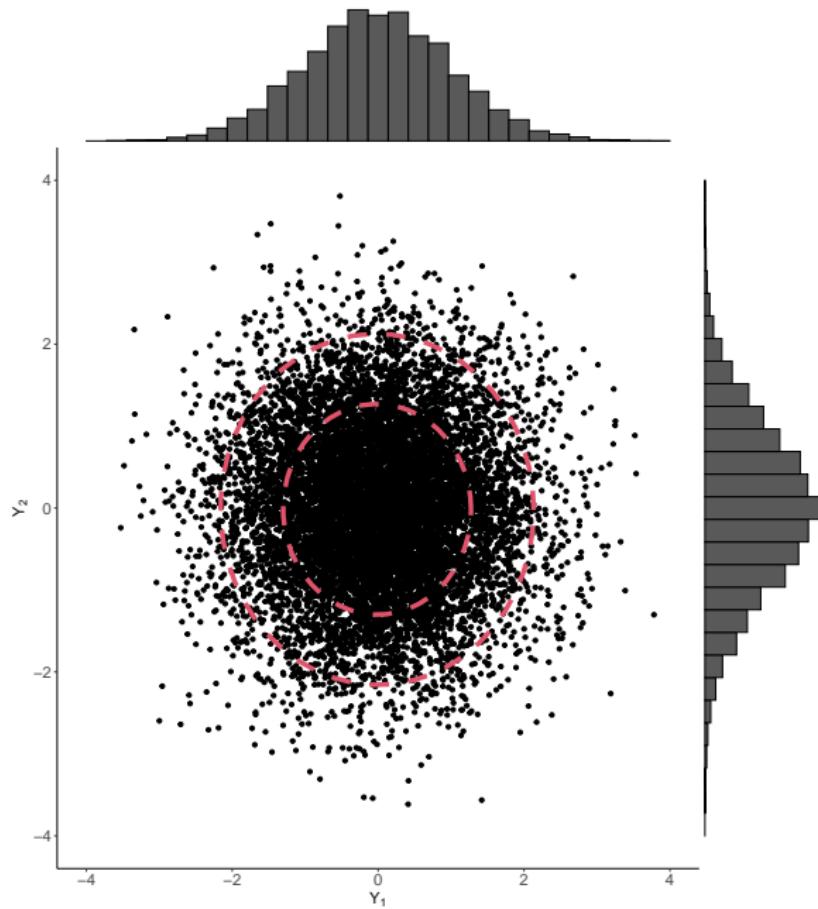
$$Y = \begin{pmatrix} Y_1 \\ Y_2 \end{pmatrix} \quad \mu = \begin{pmatrix} \mu_1 \\ \mu_2 \end{pmatrix} \quad \Sigma = \begin{pmatrix} \sigma_1^2 & \rho_{12}\sigma_1\sigma_2 \\ \rho_{21}\sigma_1\sigma_2 & \sigma_2^2 \end{pmatrix}$$

$$\text{Var}(Y_i) = \sigma_i^2 \quad \text{Cov}(Y_1, Y_2) = \rho_{12}\sigma_1\sigma_2 \quad \text{Cor}(Y_1, Y_2) = \rho_{12}$$

pdf: $f(y | \mu, \Sigma) = |\Sigma|^{-\frac{1}{2}}(2\pi)^{-\frac{d}{2}} \exp\left(-\frac{1}{2}(y - \mu)^\top \Sigma^{-1}(y - \mu)\right)$

$$\mu = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$$
$$\Sigma = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$$

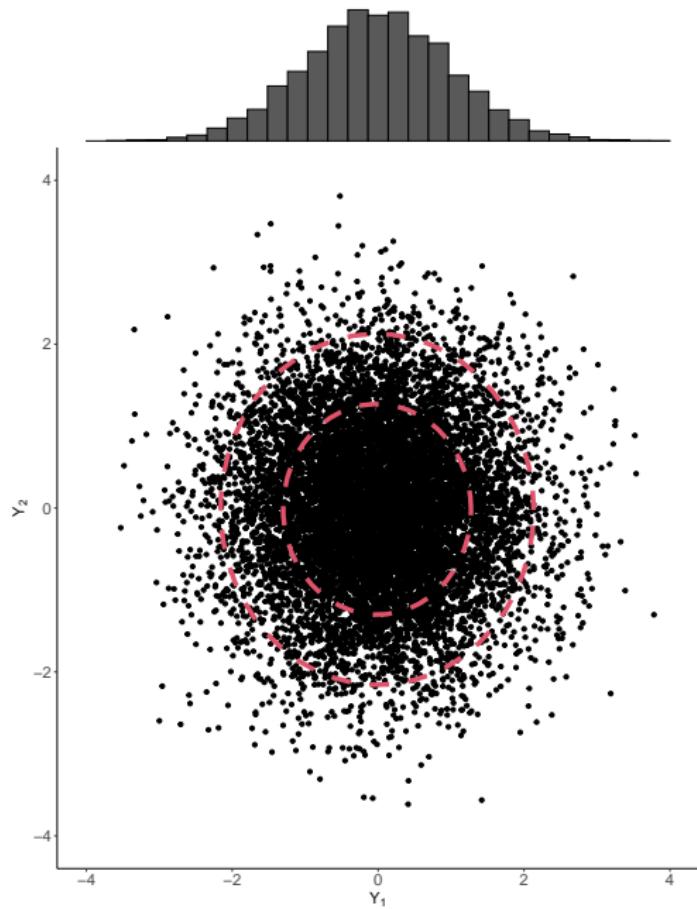




$$\mu = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$$

$$\Sigma = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$$

$\text{Cor}(Y_1, Y_2) = 0$
hence Y_1
independent of Y_2

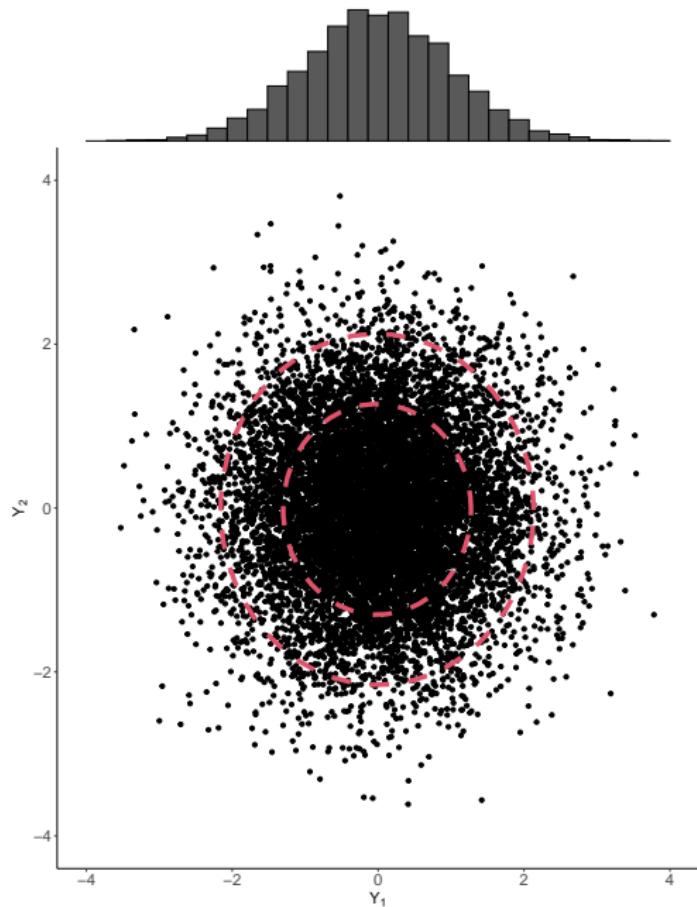


$$\mu = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$$

$$\Sigma = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$$

$\text{Cor}(Y_1, Y_2) = 0$
hence Y_1
independent of Y_2

Marginal
distributions:
 $Y_1 \sim N(0, 1)$
 $Y_2 \sim N(0, 1)$



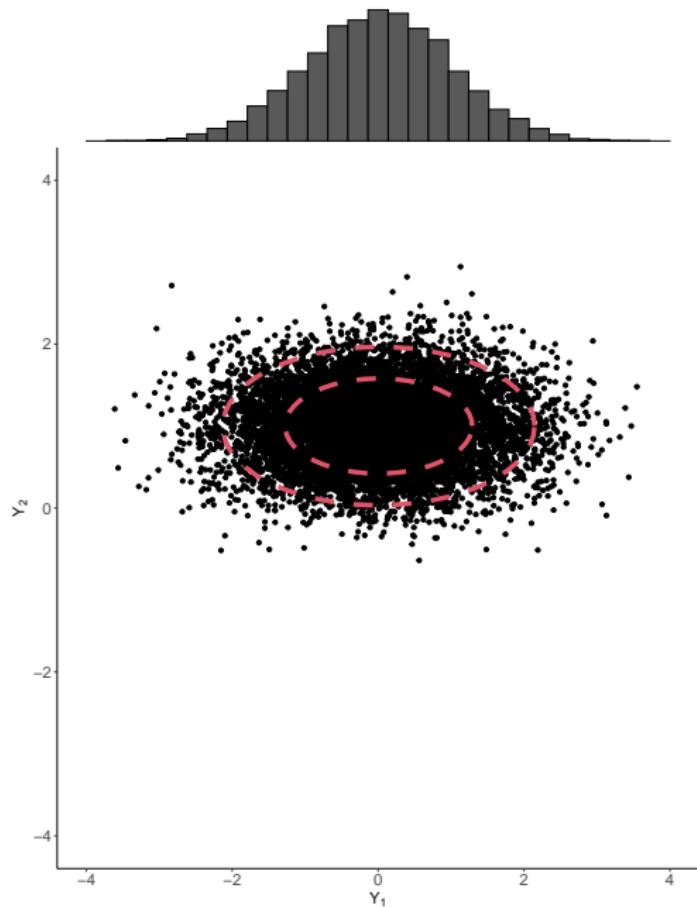
$$\mu = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$$

$$\Sigma = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$$

$\text{Cor}(Y_1, Y_2) = 0$
hence Y_1
independent of Y_2

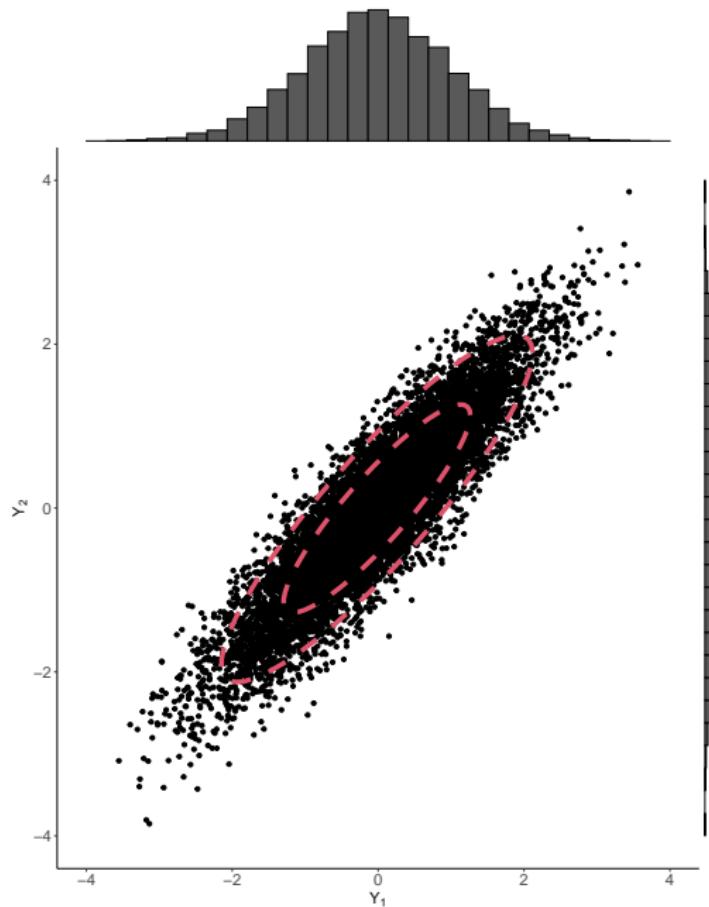
Marginal distributions:
 $Y_1 \sim N(0, 1)$
 $Y_2 \sim N(0, 1)$

Conditional distributions:
 $Y_1|Y_2 \sim N(0, 1)$
 $Y_2|Y_1 \sim N(0, 1)$

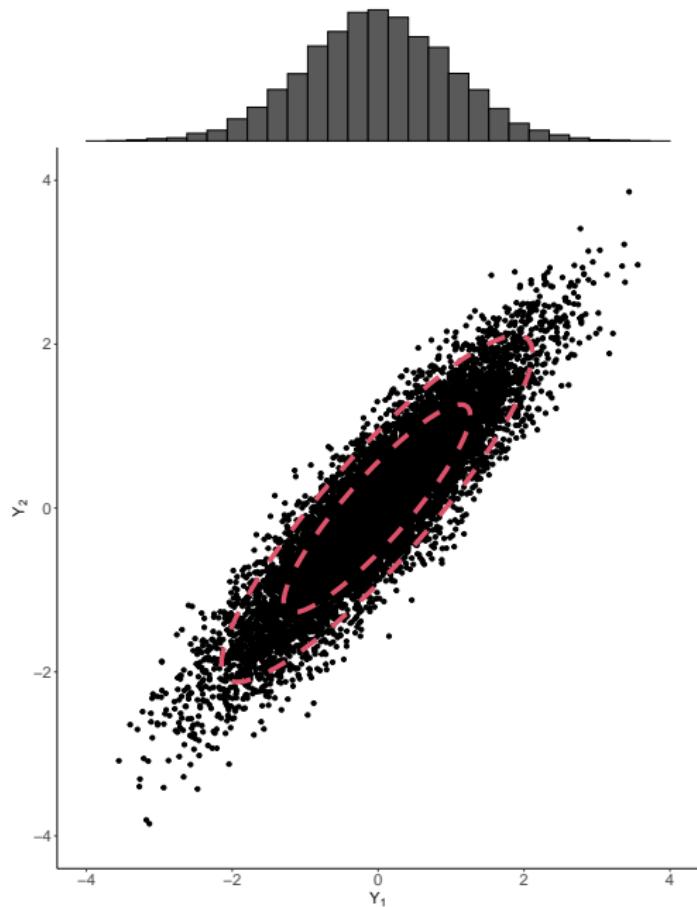


$$\mu = \begin{pmatrix} 0 \\ 1 \end{pmatrix}$$

$$\Sigma = \begin{pmatrix} 1 & 0 \\ 0 & 0.2 \end{pmatrix}$$



$$\mu = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$$
$$\Sigma = \begin{pmatrix} 1 & 0.9 \\ 0.9 & 1 \end{pmatrix}$$

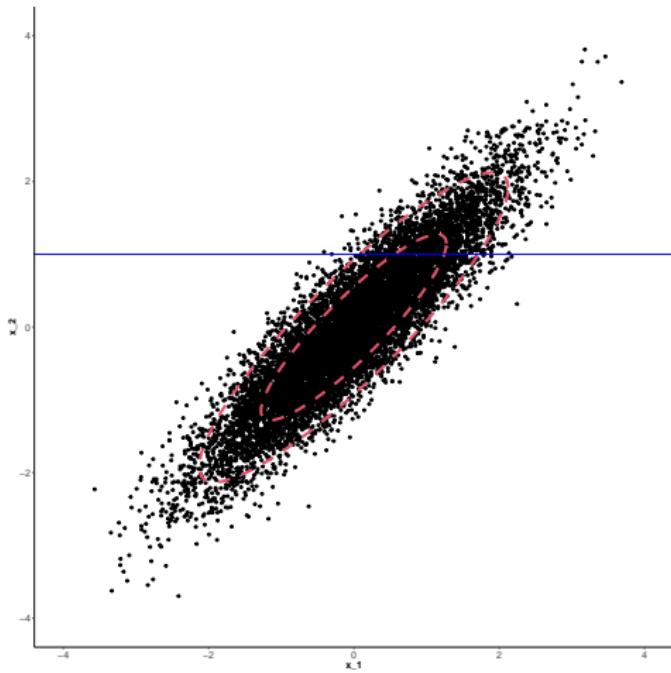


$$\mu = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$$

$$\Sigma = \begin{pmatrix} 1 & 0.9 \\ 0.9 & 1 \end{pmatrix}$$

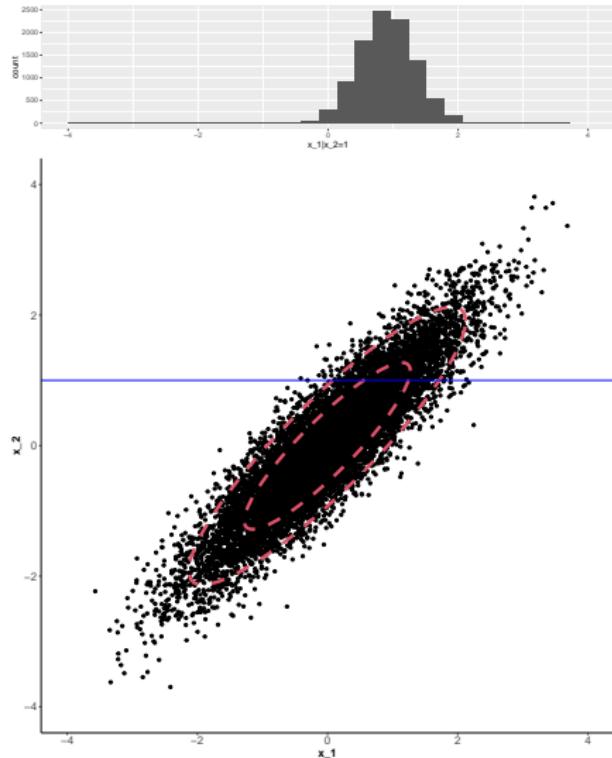
Marginal
distributions:
 $Y_1 \sim N(0, 1)$

$Y_2 \sim N(0, 1)$



$$\mu = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$$
$$\Sigma = \begin{pmatrix} 1 & 0.9 \\ 0.9 & 1 \end{pmatrix}$$

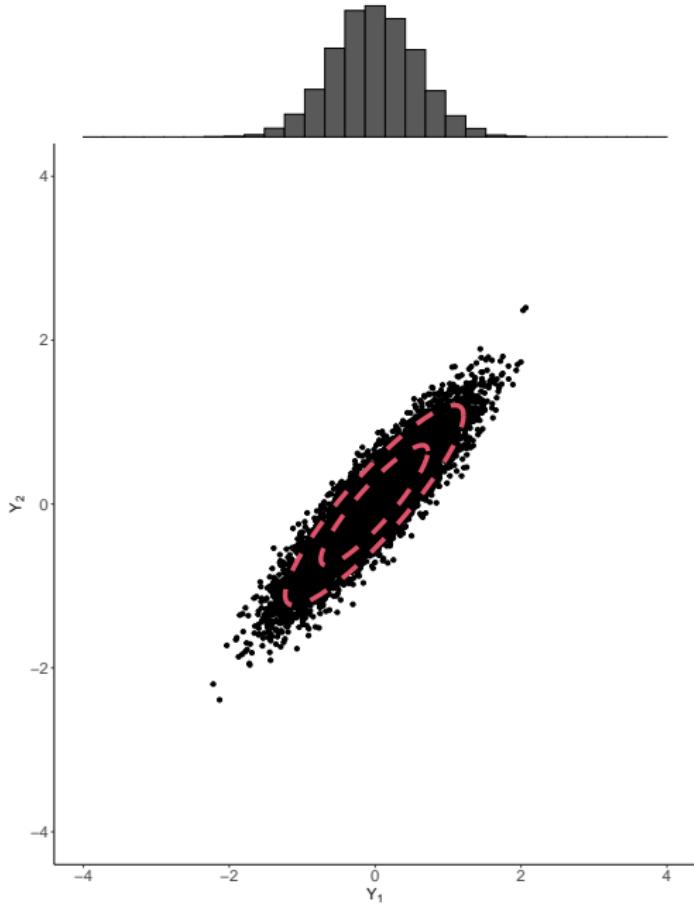
What is the
conditional
distribution of
 $Y_1|Y_2 = 1$?



$$\mu = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$$

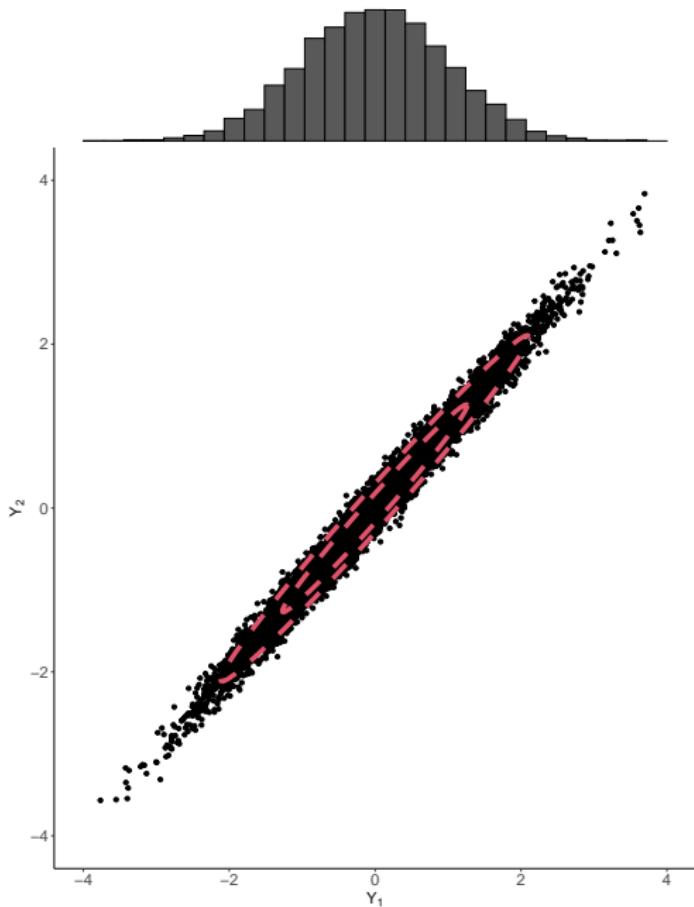
$$\Sigma = \begin{pmatrix} 1 & 0.9 \\ 0.9 & 1 \end{pmatrix}$$

$$Y_1|Y_2=1 \sim N(0.9, 0.19)$$



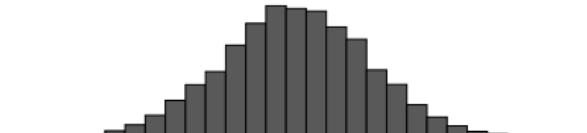
$$\mu = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$$

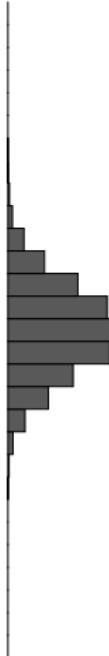
$$\Sigma = \frac{1}{3} \begin{pmatrix} 1 & 0.9 \\ 0.9 & 1 \end{pmatrix}$$

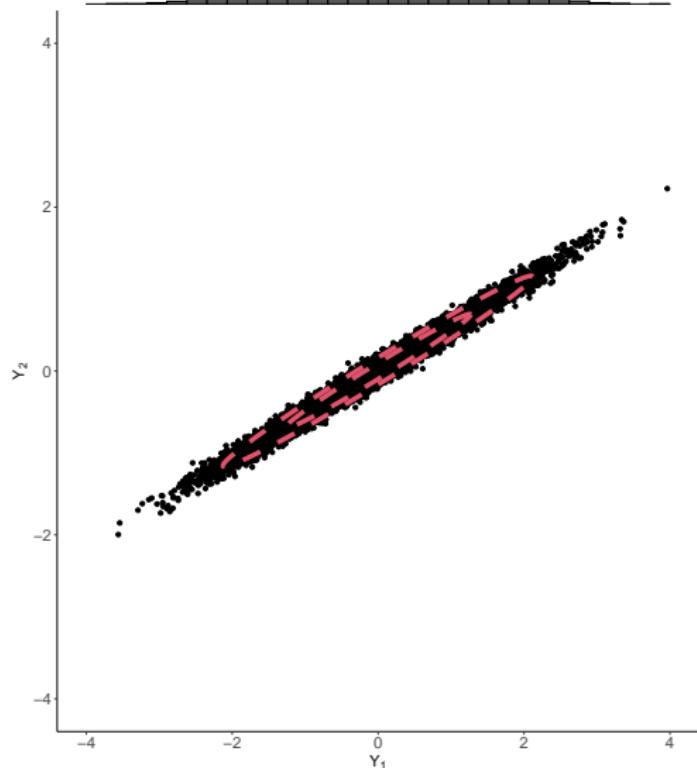


$$\mu = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$$

$$\Sigma = \begin{pmatrix} 1 & 0.99 \\ 0.99 & 1 \end{pmatrix}$$


$$\mu = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$$


$$\Sigma = \begin{pmatrix} 1 & 0.54 \\ 0.54 & 0.3 \end{pmatrix}$$

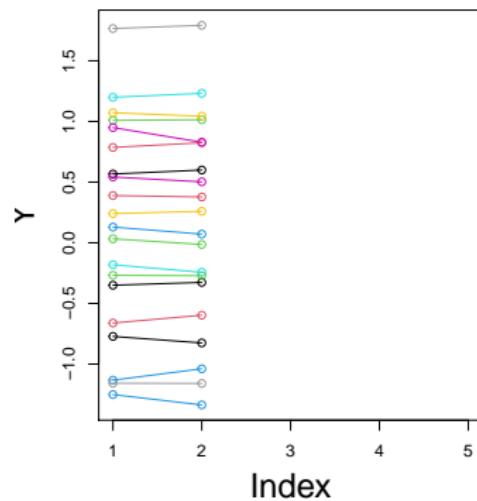
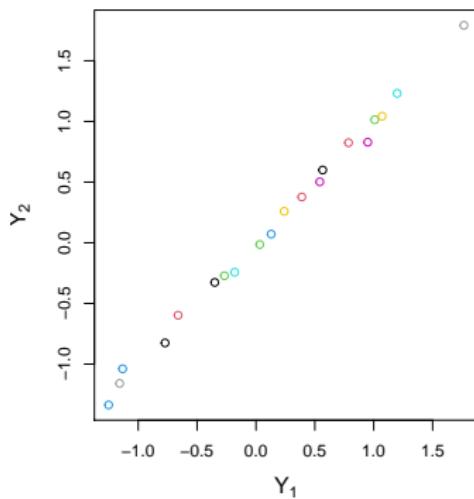

$$Cor(Y_1, Y_2) = \\ 0.54 / \sqrt(0.3) = \\ 0.99$$

More pictures

Hard to visualise in dimensions > 2, so stack points next to each other.

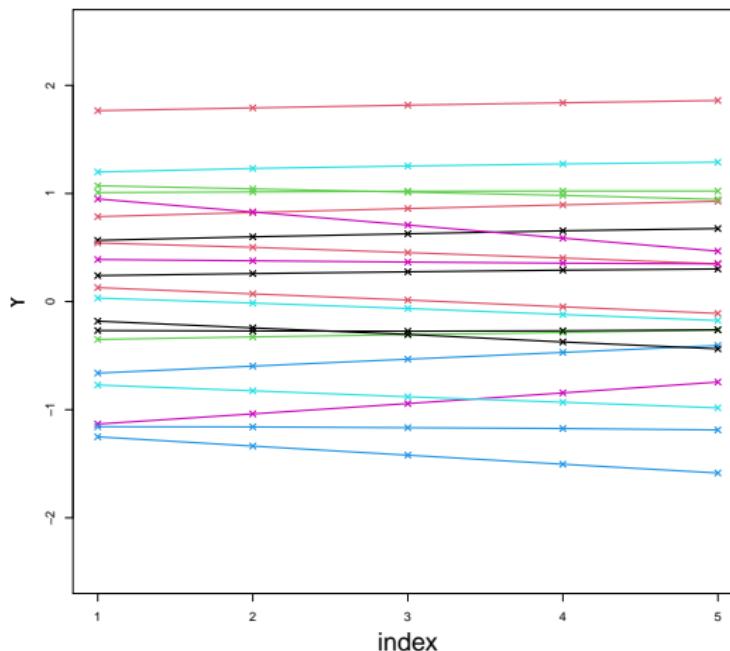
More pictures

Hard to visualise in dimensions > 2, so stack points next to each other.
So for 2d instead of we have



Consider $d = 5$ with

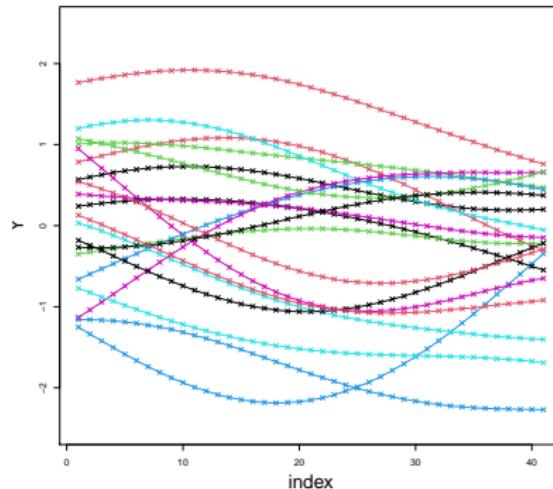
$$\mu = \begin{pmatrix} 0 \\ 0 \\ \vdots \\ 0 \end{pmatrix} \quad \Sigma = \begin{pmatrix} 1 & 0.99 & 0.98 & 0.97 & 0.96 \\ 0.99 & 1 & 0.99 & 0.98 & 0.97 \\ 0.98 & 0.99 & 1 & 0.99 & 0.98 \\ 0.97 & 0.98 & 0.99 & 1 & 0.99 \\ 0.96 & 0.97 & 0.98 & 0.99 & 1 \end{pmatrix}$$



Each line is one sample.

$$d = 50$$

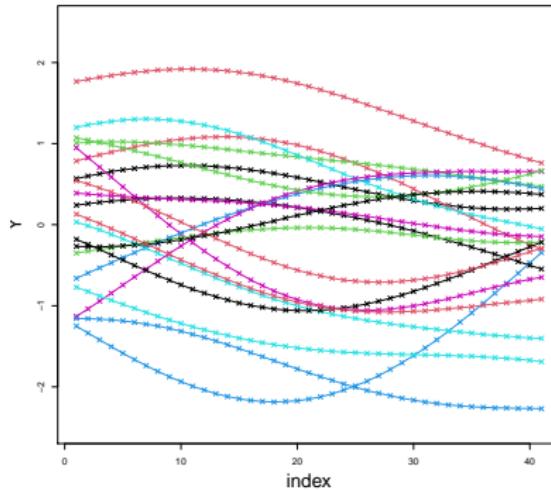
$$\mu = \begin{pmatrix} 0 \\ 0 \\ \vdots \\ 0 \end{pmatrix} \quad \Sigma = \begin{pmatrix} 1 & 0.99 & 0.98 & 0.97 & 0.96 & \dots \\ 0.99 & 1 & 0.99 & 0.98 & 0.97 & \dots \\ 0.98 & 0.99 & 1 & 0.99 & 0.98 & \dots \\ 0.97 & 0.98 & 0.99 & 1 & 0.99 & \dots \\ 0.96 & 0.97 & 0.98 & 0.99 & 1 & \dots \\ \vdots & \vdots & \vdots & \vdots & \vdots & \ddots \end{pmatrix}$$



Each line is one sample.

$$d = 50$$

$$\mu = \begin{pmatrix} 0 \\ 0 \\ \vdots \\ 0 \end{pmatrix} \quad \Sigma = \begin{pmatrix} 1 & 0.99 & 0.98 & 0.97 & 0.96 & \dots \\ 0.99 & 1 & 0.99 & 0.98 & 0.97 & \dots \\ 0.98 & 0.99 & 1 & 0.99 & 0.98 & \dots \\ 0.97 & 0.98 & 0.99 & 1 & 0.99 & \dots \\ 0.96 & 0.97 & 0.98 & 0.99 & 1 & \dots \\ \vdots & \vdots & \vdots & \vdots & \vdots & \ddots \end{pmatrix}$$



Each line is one sample.

We can think of Gaussian processes as an infinite dimensional distribution over functions - all we need to do is change the indexing

- We can let the index be x and let x take values in \mathbb{R} .

Gaussian processes

A stochastic process is a collection of random variables indexed by some variable $x \in \mathcal{X}$

$$y = \{y(x) : x \in \mathcal{X}\}$$

Gaussian processes

A stochastic process is a collection of random variables indexed by some variable $x \in \mathcal{X}$

$$y = \{y(x) : x \in \mathcal{X}\}$$

Usually $y(x) \in \mathbb{R}$ and $\mathcal{X} \subset \mathbb{R}^n$

Gaussian processes

A stochastic process is a collection of random variables indexed by some variable $x \in \mathcal{X}$

$$y = \{y(x) : x \in \mathcal{X}\}$$

Usually $y(x) \in \mathbb{R}$ and $\mathcal{X} \subset \mathbb{R}^n$ - think of y as a function of x .

Gaussian processes

A stochastic process is a collection of random variables indexed by some variable $x \in \mathcal{X}$

$$y = \{y(x) : x \in \mathcal{X}\}$$

Usually $y(x) \in \mathbb{R}$ and $\mathcal{X} \subset \mathbb{R}^n$ - think of y as a function of x .

If $|\mathcal{X}| = \infty$, y is an infinite dimensional process, e.g., $|\mathcal{X}| = [0, 1]$

Gaussian processes

A stochastic process is a collection of random variables indexed by some variable $x \in \mathcal{X}$

$$y = \{y(x) : x \in \mathcal{X}\}$$

Usually $y(x) \in \mathbb{R}$ and $\mathcal{X} \subset \mathbb{R}^n$ - think of y as a function of x .

If $|\mathcal{X}| = \infty$, y is an infinite dimensional process, e.g., $|\mathcal{X}| = [0, 1]$

Thankfully, the law of y is uniquely determined by the finite dimensional distributions (FDDs), i.e., for all x_1, \dots, x_n and for all $n \in \mathbb{N}$

$$\mathbb{P}(y(x_1) \leq c_1, \dots, y(x_n) \leq c_n)$$

Gaussian processes

A stochastic process is a collection of random variables indexed by some variable $x \in \mathcal{X}$

$$y = \{y(x) : x \in \mathcal{X}\}$$

Usually $y(x) \in \mathbb{R}$ and $\mathcal{X} \subset \mathbb{R}^n$ - think of y as a function of x .

If $|\mathcal{X}| = \infty$, y is an infinite dimensional process, e.g., $|\mathcal{X}| = [0, 1]$

Thankfully, the law of y is uniquely determined by the finite dimensional distributions (FDDs), i.e., for all x_1, \dots, x_n and for all $n \in \mathbb{N}$

$$\mathbb{P}(y(x_1) \leq c_1, \dots, y(x_n) \leq c_n)$$

A **Gaussian process** is a stochastic process with Gaussian FDDs, i.e.,

$$(y(x_1), \dots, y(x_n)) \sim N_n(\mu, \Sigma)$$

Gaussian processes

A stochastic process is a collection of random variables indexed by some variable $x \in \mathcal{X}$

$$y = \{y(x) : x \in \mathcal{X}\}$$

Usually $y(x) \in \mathbb{R}$ and $\mathcal{X} \subset \mathbb{R}^n$ - think of y as a function of x .

If $|\mathcal{X}| = \infty$, y is an infinite dimensional process, e.g., $|\mathcal{X}| = [0, 1]$

Thankfully, the law of y is uniquely determined by the finite dimensional distributions (FDDs), i.e., for all x_1, \dots, x_n and for all $n \in \mathbb{N}$

$$\mathbb{P}(y(x_1) \leq c_1, \dots, y(x_n) \leq c_n)$$

A **Gaussian process** is a stochastic process with Gaussian FDDs, i.e.,

$$(y(x_1), \dots, y(x_n)) \sim N_n(\mu, \Sigma)$$

Write $y(\cdot) \sim GP$ to denote that we model the *function* y as a GP.

Mean and covariance function

To fully specify the law of a Gaussian *distribution* we only need the mean and variance.

$$X \sim N(\mu, \Sigma)$$

Mean and covariance function

To fully specify the law of a Gaussian *distribution* we only need the mean and variance.

$$X \sim N(\mu, \Sigma)$$

To fully specify the law of a Gaussian *process*, we need to specify mean and covariance **functions**.

Mean and covariance function

To fully specify the law of a Gaussian *distribution* we only need the mean and variance.

$$X \sim N(\mu, \Sigma)$$

To fully specify the law of a Gaussian *process*, we need to specify mean and covariance **functions**.

$$y(\cdot) \sim GP(m(\cdot), k(\cdot, \cdot))$$

Mean and covariance function

To fully specify the law of a Gaussian *distribution* we only need the mean and variance.

$$X \sim N(\mu, \Sigma)$$

To fully specify the law of a Gaussian *process*, we need to specify mean and covariance **functions**.

$$y(\cdot) \sim GP(m(\cdot), k(\cdot, \cdot))$$

where

$$\mathbb{E}(y(x)) = m(x)$$

$$\text{Cov}(y(x), y(x')) = k(x, x')$$

Mean and covariance function

To fully specify the law of a Gaussian *distribution* we only need the mean and variance.

$$X \sim N(\mu, \Sigma)$$

To fully specify the law of a Gaussian *process*, we need to specify mean and covariance **functions**.

$$y(\cdot) \sim GP(m(\cdot), k(\cdot, \cdot))$$

where

$$\begin{aligned}\mathbb{E}(y(x)) &= m(x) \\ \text{Cov}(y(x), y(x')) &= k(x, x')\end{aligned}$$

In GP software, these are often the only choices we need to make.

Specifying the mean function

We are free to choose the mean $m(x) = \mathbb{E}(y(x))$ and covariance $k(x, x') = \text{Cov}(y(x), y(x'))$ functions however we like (e.g. trial and error), subject to some 'rules':

Specifying the mean function

We are free to choose the mean $m(x) = \mathbb{E}(y(x))$ and covariance $k(x, x') = \text{Cov}(y(x), y(x'))$ functions however we like (e.g. trial and error), subject to some 'rules':

- We can use any mean function we want:

$$m(x) = \mathbb{E}(y(x))$$

Most popular choices are $m(x) = 0$ or $m(x) = \text{const}$ for all x , or $m(x) = \beta^\top x$

Covariance functions

We usually use a covariance function that is a function of the indexes/locations

$$k(x, x') = \text{Cov}(y(x), y(x')),$$

Covariance functions

We usually use a covariance function that is a function of the indexes/locations

$$k(x, x') = \text{Cov}(y(x), y(x')),$$

k must be a **positive semi-definite function**, i.e., lead to valid covariance matrices:

- Given locations x_1, \dots, x_n , the $n \times n$ Gram matrix K with
 $K_{ij} = k(x_i, x_j)$

$$\begin{pmatrix} k(x_1, x_1) & k(x_1, x_2) & \dots & k(x_1, x_n) \\ \vdots & & & \vdots \\ k(x_n, x_1) & k(x_n, x_2) & \dots & k(x_n, x_n) \end{pmatrix}$$

must be a positive semi-definite matrix (ie a covariance matrix).

Covariance functions

We usually use a covariance function that is a function of the indexes/locations

$$k(x, x') = \text{Cov}(y(x), y(x')),$$

k must be a **positive semi-definite function**, i.e., lead to valid covariance matrices:

- Given locations x_1, \dots, x_n , the $n \times n$ Gram matrix K with
 $K_{ij} = k(x_i, x_j)$

$$\begin{pmatrix} k(x_1, x_1) & k(x_1, x_2) & \dots & k(x_1, x_n) \\ \vdots & & & \vdots \\ k(x_n, x_1) & k(x_n, x_2) & \dots & k(x_n, x_n) \end{pmatrix}$$

must be a positive semi-definite matrix (ie a covariance matrix).

- This can be problematic...
 - Difficult to create semi-definite functions $k(x, x')$

We often assume k is a function of only the distance between locations

$$\text{Cov}(y(x), y(x')) = k(x - x')$$

which results in a **stationary** process.

We often assume k is a function of only the distance between locations

$$\text{Cov}(y(x), y(x')) = k(x - x')$$

which results in a **stationary** process.

If $\text{Cov}(y(x), y(x')) = k(||x - x'||)$ the covariance function is said to be **isotropic**.

We often assume k is a function of only the distance between locations

$$\text{Cov}(y(x), y(x')) = k(x - x')$$

which results in a **stationary** process.

If $\text{Cov}(y(x), y(x')) = k(||x - x'||)$ the covariance function is said to be **isotropic**.

The covariance function determines the *nature* of the GP.

- k determines the hypothesis space/space of functions

We often assume k is a function of only the distance between locations

$$\text{Cov}(y(x), y(x')) = k(x - x')$$

which results in a **stationary** process.

If $\text{Cov}(y(x), y(x')) = k(\|x - x'\|)$ the covariance function is said to be **isotropic**.

The covariance function determines the *nature* of the GP.

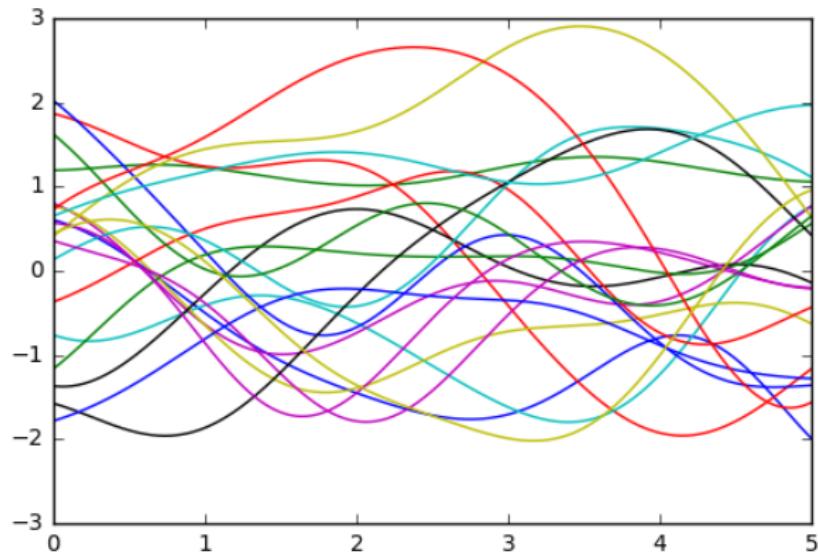
- k determines the hypothesis space/space of functions

We usually build k by selecting from a small candidate set of functions (that contain free parameters), and modifying them according to some rules.

Examples

RBF/Squared-exponential/exponentiated quadratic

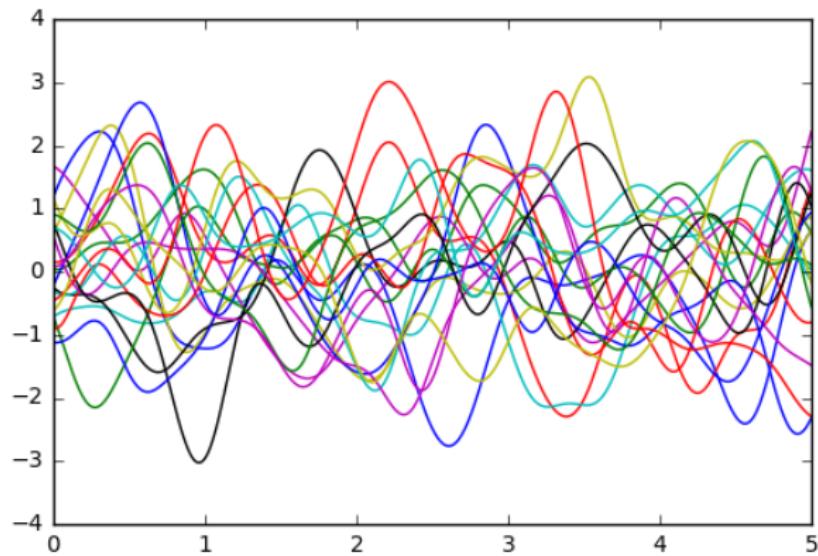
$$k(x, x') = \exp\left(-\frac{1}{2}(x - x')^2\right)$$



Examples

RBF/Squared-exponential/exponentiated quadratic

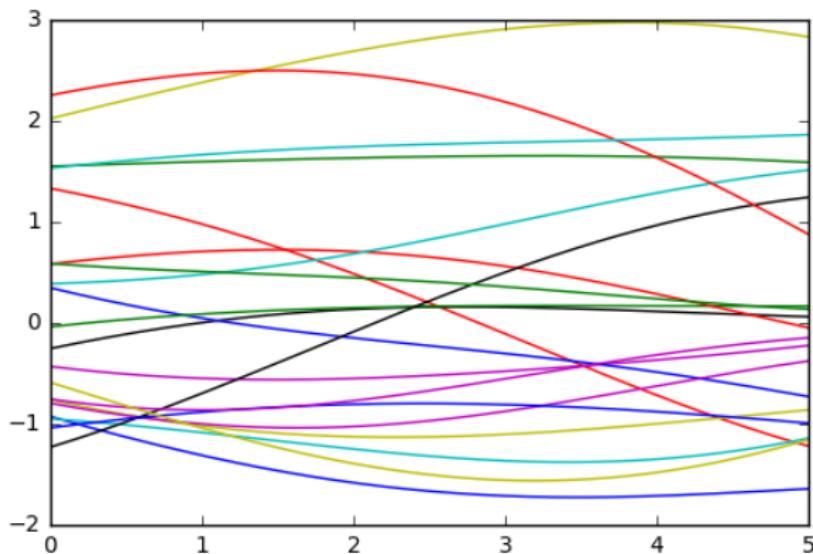
$$k(x, x') = \exp\left(-\frac{1}{2} \frac{(x - x')^2}{0.25^2}\right)$$



Examples

RBF/Squared-exponential/exponentiated quadratic

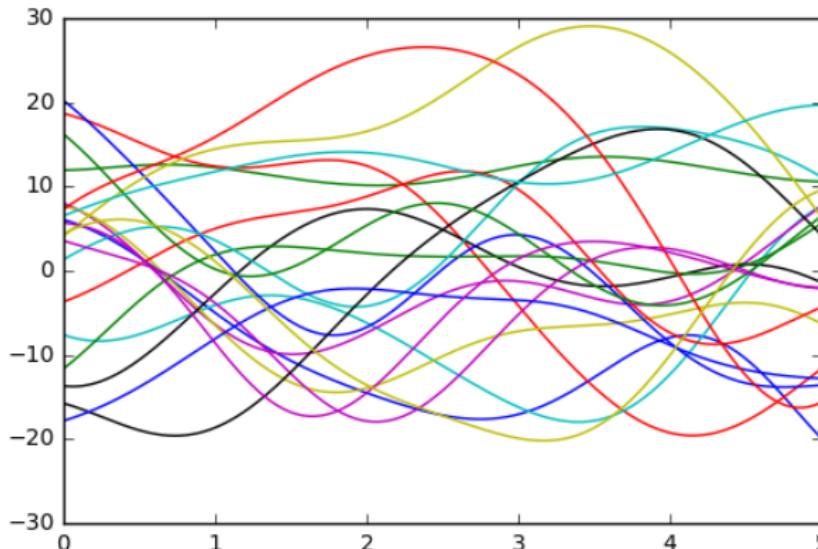
$$k(x, x') = \exp\left(-\frac{1}{2} \frac{(x - x')^2}{4^2}\right)$$



Examples

RBF/Squared-exponential/exponentiated quadratic

$$k(x, x') = \textcolor{red}{100} \exp\left(-\frac{1}{2}(x - x')^2\right)$$



$$\text{In general } k(x, x') = \sigma^2 \exp\left(-\frac{1}{2} \frac{(x - x')^2}{\lambda^2}\right)$$

Examples

Matérn Covariance functions

The RBF kernel produces functions that are very smooth (infinitely differentiable)

The Matérn family of covariance functions allows us to control the degree of differentiability via an additional parameter ν

Examples

Matérn Covariance functions

The RBF kernel produces functions that are very smooth (infinitely differentiable)

The Matérn family of covariance functions allows us to control the degree of differentiability via an additional parameter ν

- $\nu = 1/2$: $k(x, x') = \sigma^2 \exp\left(-\frac{|x-x'|}{\lambda}\right)$ gives functions that are no-where differentiable

Examples

Matérn Covariance functions

The RBF kernel produces functions that are very smooth (infinitely differentiable)

The Matérn family of covariance functions allows us to control the degree of differentiability via an additional parameter ν

- $\nu = 1/2$: $k(x, x') = \sigma^2 \exp\left(-\frac{|x-x'|}{\lambda}\right)$ gives functions that are no-where differentiable
- $\nu = 3/2$: $k(x, x') = \sigma^2 \left(1 + \frac{\sqrt{3}|x-x'|}{\lambda}\right) \exp\left(-\frac{\sqrt{3}|x-x'|}{\lambda}\right)$ gives functions that are once differentiable

Examples

Matérn Covariance functions

The RBF kernel produces functions that are very smooth (infinitely differentiable)

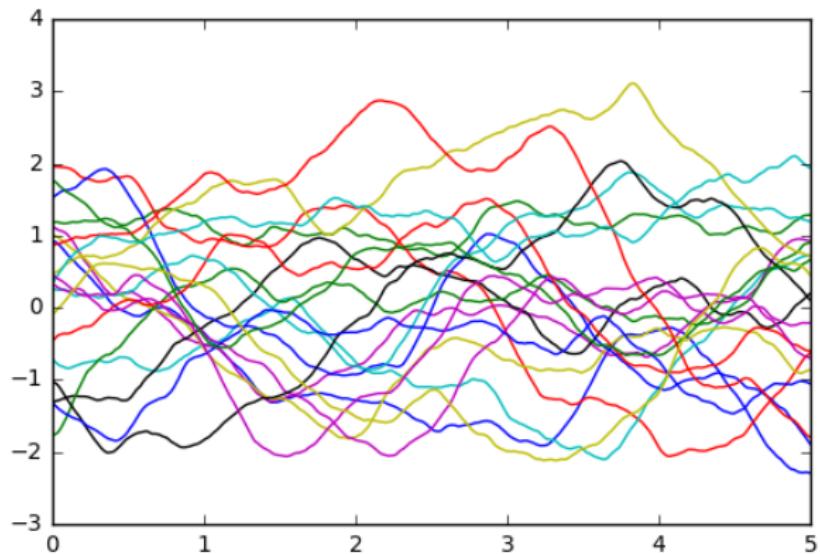
The Matérn family of covariance functions allows us to control the degree of differentiability via an additional parameter ν

- $\nu = 1/2$: $k(x, x') = \sigma^2 \exp\left(-\frac{|x-x'|}{\lambda}\right)$ gives functions that are no-where differentiable
- $\nu = 3/2$: $k(x, x') = \sigma^2 \left(1 + \frac{\sqrt{3}|x-x'|}{\lambda}\right) \exp\left(-\frac{\sqrt{3}|x-x'|}{\lambda}\right)$ gives functions that are once differentiable
- $\nu = 5/2$: $k(x, x') = \sigma^2 \left(1 + \frac{\sqrt{5}|x-x'|}{\lambda} + \frac{5|x-x'|^2}{3\lambda^2}\right) \exp\left(-\frac{\sqrt{5}|x-x'|}{\lambda}\right)$ gives functions that are twice differentiable

Examples

Matern 3/2

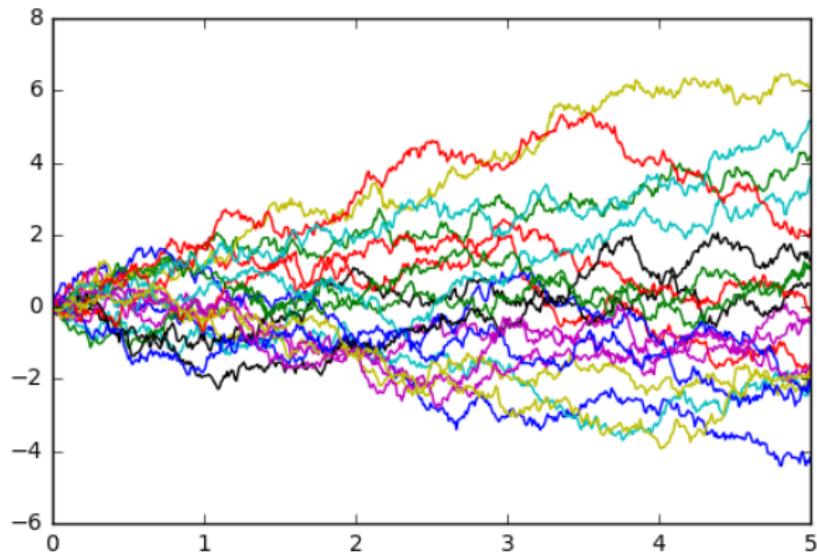
$$k(x, x') \sim (1 + |x - x'|) \exp(-|x - x'|)$$



Examples

Brownian motion

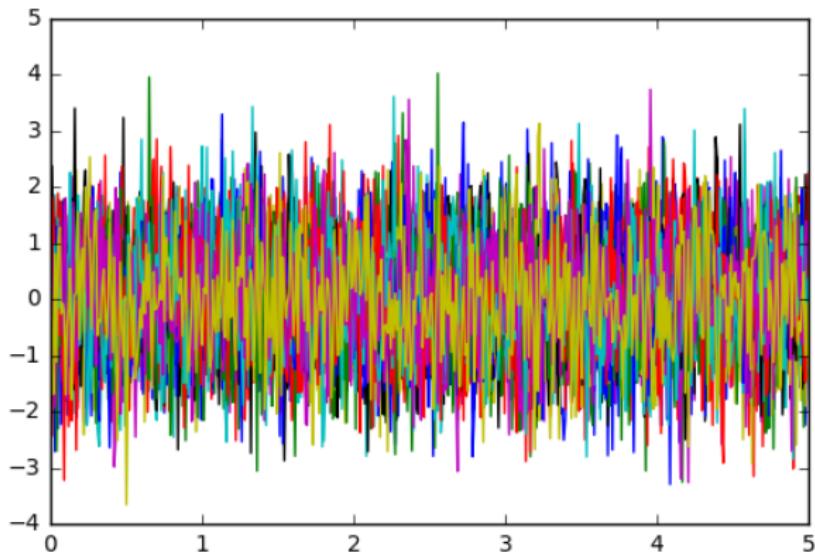
$$k(x, x') = \min(x, x')$$



Examples

White noise

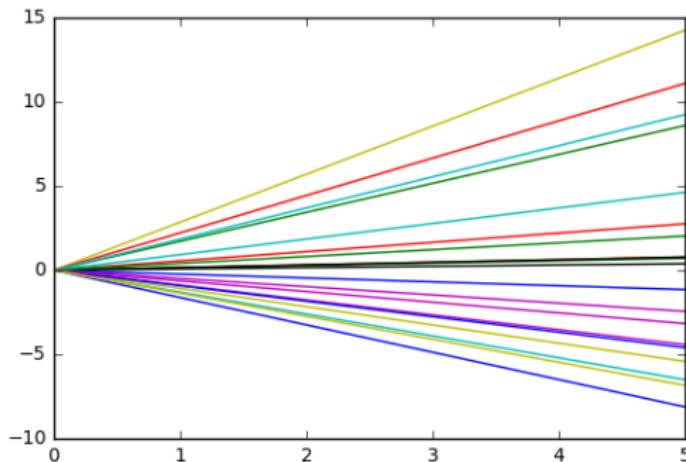
$$k(x, x') = \begin{cases} 1 & \text{if } x = x' \\ 0 & \text{otherwise} \end{cases}$$



Examples

A final example:

$$k(x, x') = xx'$$

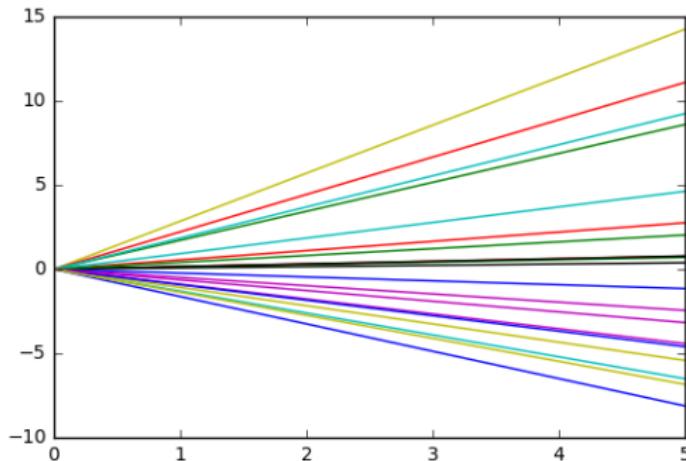


What is happening?

Examples

A final example:

$$k(x, x') = xx'$$



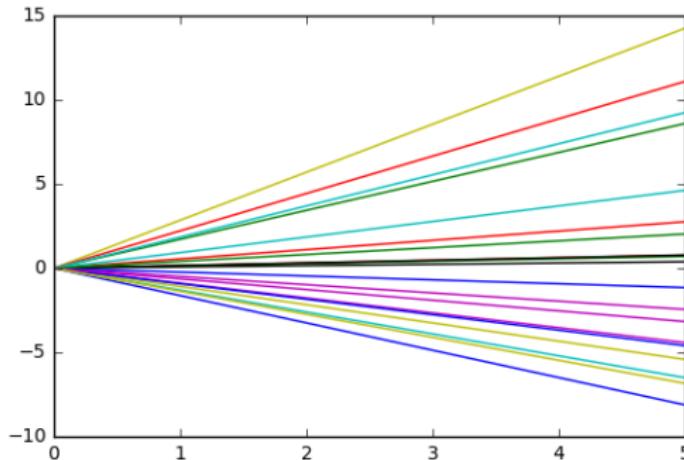
What is happening?

Suppose $y(x) = cx$ where $c \sim N(0, 1)$.

Examples

A final example:

$$k(x, x') = xx'$$



What is happening?

Suppose $y(x) = cx$ where $c \sim N(0, 1)$.

Then

$$\begin{aligned}\text{Cov}(y(x), y(x')) &= \text{Cov}(cx, cx') = x\text{Cov}(c, c)x' \\ &= xx'\end{aligned}$$

So $y(\cdot) \sim GP(0, k(x, x'))$ with $k(x, x') = xx'$

Rules for combining covariance functions

Let k_1 and k_2 be valid covariance functions. Then

- $k_1 + k_2$ is a valid covariance function (even if k_1 and k_2 take different inputs).
- $k_1 k_2$ is a valid covariance function (also true if different inputs).
- For any function g , $k_1(g(x), g(x'))$ is a valid covariance function as is $g(x)k_1(x, x')g(x')$

Using these rules we can combine our standard set of positive definite covariance functions to create a richer family of covariance functions.

Choosing kernels and hyperparameters

GP properties are inherited primarily from the covariance function k .

- Continuity
- Differentiability
- Variance and length-scale

2f is mean square cts at x if for all sequences $x_k \rightarrow x$ we have $\mathbb{E}(f(x_k) - f(x))^2 \rightarrow 0$

Choosing kernels and hyperparameters

GP properties are inherited primarily from the covariance function k .

- Continuity
 - ▶ $f(x) \sim GP$ is (mean square²) continuous at x^* if $k(x, x')$ and $m(x)$ are continuous at $x = x' = x^*$
 - ▶ For stationary kernels, only require continuity at $k(0)$

- Differentiability
 - ▶ $f(x) \sim GP$ is (mean square) differentiable if $k'(x, x') = \frac{\partial^2}{\partial x \partial x'} k(x, x')$ exists.

- Variance and length-scale

² f is mean square cts at x if for all sequences $x_k \rightarrow x$ we have $\mathbb{E}(f(x_k) - f(x))^2 \rightarrow 0$

Choosing kernels and hyperparameters

GP properties are inherited primarily from the covariance function k .

- Continuity
 - ▶ $f(x) \sim GP$ is (mean square²) continuous at x^* if $F_k(x, x')$ and $m(x)$ are continuous at $x = x' = x^*$
 - ▶ For stationary kernels, only require continuity at $k(0)$

- Differentiability
 - ▶ $f(x) \sim GP$ is (mean square) differentiable if $k'(x, x') = \frac{\partial^2}{\partial x \partial x'} k(x, x')$ exists.

- Variance and length-scale controlled by hyper-parameters $k = k_\psi$:
 - ▶ how much f varies between samples
 - ▶ how fast $f(x)$ changes with x within a sample.

² f is mean square cts at x if for all sequences $x_k \rightarrow x$ we have $\mathbb{E}(f(x_k) - f(x))^2 \rightarrow 0$

Choosing kernels and hyperparameters

GP properties are inherited primarily from the covariance function k .

- Continuity
 - ▶ $f(x) \sim GP$ is (mean square²) continuous at x^* if $k(x, x')$ and $m(x)$ are continuous at $x = x' = x^*$
 - ▶ For stationary kernels, only require continuity at $k(0)$

- Differentiability
 - ▶ $f(x) \sim GP$ is (mean square) differentiable if $k'(x, x') = \frac{\partial^2}{\partial x \partial x'} k(x, x')$ exists.

- Variance and length-scale controlled by hyper-parameters $k = k_\psi$:
 - ▶ how much f varies between samples
 - ▶ how fast $f(x)$ changes with x within a sample.

Typically choose the family of kernels by

- measures of fit (marginal likelihood, Bayes factors, ...)
- predictive skill (held-out data, cross-validation, ...)

Choose hyperparameters by maximum likelihood, Bayes, etc.

² f is mean square cts at x if for all sequences $x_k \rightarrow x$ we have $\mathbb{E}(f(x_k) - f(x))^2 \rightarrow 0$

Conditional updates of Gaussian processes

If f is a Gaussian process $f(\cdot) \sim GP(m(\cdot), k(\cdot, \cdot))$, then

$$f(x_1), \dots, f(x_n), f(x) \sim N_{n+1}(\mu, \Sigma)$$

where

Conditional updates of Gaussian processes

If f is a Gaussian process $f(\cdot) \sim GP(m(\cdot), k(\cdot, \cdot))$, then

$$f(x_1), \dots, f(x_n), f(x) \sim N_{n+1}(\mu, \Sigma)$$

where

$$\mu = \begin{pmatrix} m(x_1) \\ \vdots \\ m(x_n) \\ \hline m(x) \end{pmatrix} = \begin{pmatrix} | \\ m_X \\ | \\ m(x) \end{pmatrix}$$

and

$$\Sigma = \left(\begin{array}{ccc|c} k(x_1, x_1) & \dots & k(x_1, x_n) & k(x_1, x) \\ \vdots & & \vdots & \vdots \\ k(x_n, x_1) & \dots & k(x_n, x_n) & k(x_n, x) \\ \hline k(x, x_1) & \dots & k(x, x_n) & k(x, x) \end{array} \right) = \left(\begin{array}{c|c} K_{XX} & k_X(x) \\ \hline k_X(x)^T & k(x, x) \end{array} \right)$$

where $X = \{x_1, \dots, x_n\}$, $[K_{XX}]_{ij} = k(x_i, x_j)$ is the Gram/kernel matrix,

Conditional updates of Gaussian processes

Then

$$f(x) | f(x_1), \dots, f(x_n) \sim N(\bar{m}(x), \bar{k}(x))$$

Conditional updates of Gaussian processes

Then

$$f(x) | f(x_1), \dots, f(x_n) \sim N(\bar{m}(x), \bar{k}(x))$$

where

$$\bar{m}(x) = m(x) + k_x(x)^\top K_{xx}^{-1} (\mathbf{f} - m_x)$$

$$\bar{k}(x) = k(x, x) - k_x(x)^\top K_{xx}^{-1} k_x(x)$$

where

$$\mathbf{f} = (f(x_1), \dots, f(x_n))^\top$$

Conditional updates of Gaussian processes

Then

$$f(x) | f(x_1), \dots, f(x_n) \sim N(\bar{m}(x), \bar{k}(x))$$

where

$$\bar{m}(x) = m(x) + k_x(x)^\top K_{xx}^{-1} (\mathbf{f} - mx)$$

$$\bar{k}(x) = k(x, x) - k_x(x)^\top K_{xx}^{-1} k_x(x)$$

where

$$\mathbf{f} = (f(x_1), \dots, f(x_n))^\top$$

The **prior** mean and covariance functions m and k , are replaced by the **posterior** mean and covariance functions \bar{m} and \bar{k} ,

Summary: Conditional updates of Gaussian processes

If f is a Gaussian process,

$$f(\cdot) \sim GP(m(\cdot), k(\cdot, \cdot))$$

then if we observe f at x_1, \dots, x_n , then

$$f(\cdot) | f(x_1), \dots, f(x_n) \sim GP(\bar{m}(\cdot), \bar{k}(\cdot, \cdot))$$

where \bar{m} and \bar{k} are the *posterior* mean and covariance functions on the previous slide.

Summary: Conditional updates of Gaussian processes

If f is a Gaussian process,

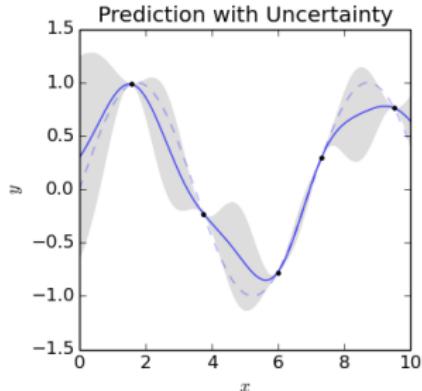
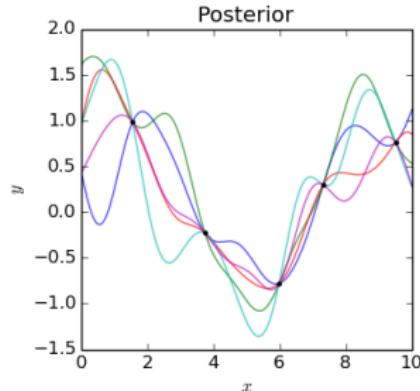
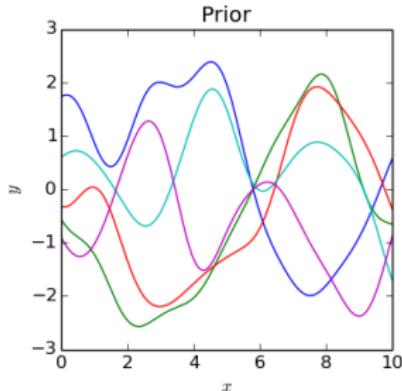
$$f(\cdot) \sim GP(m(\cdot), k(\cdot, \cdot))$$

then if we observe f at x_1, \dots, x_n , then

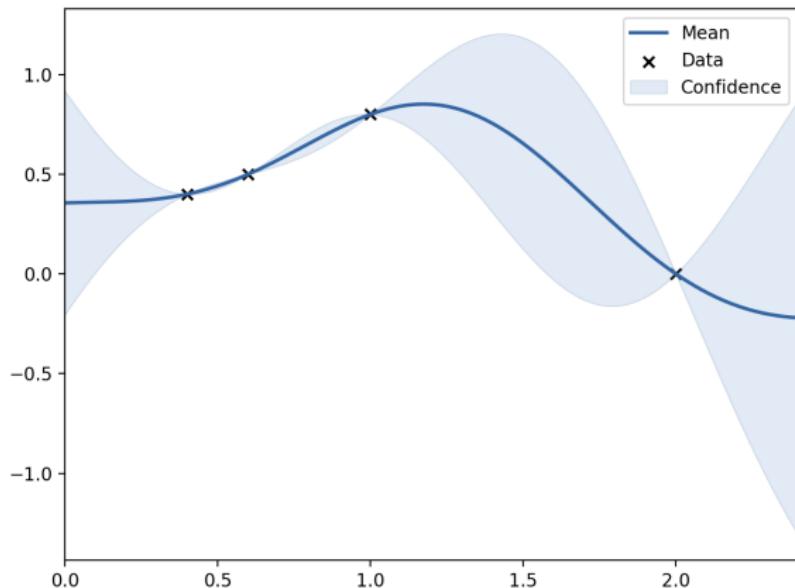
$$f(\cdot) | f(x_1), \dots, f(x_n) \sim GP(\bar{m}(\cdot), \bar{k}(\cdot, \cdot))$$

where \bar{m} and \bar{k} are the *posterior* mean and covariance functions on the previous slide.

- f is still a GP even though we've observed its value at a number of locations.



No noise/nugget - Interpolation



Solid line $\bar{m}(x) = m(x) + k_X(x)^\top K_{XX}^{-1}(\mathbf{f} - m_X)$

Shaded region $\bar{m}(x) \pm 1.96\sqrt{\bar{k}(x)}$

$$\bar{k}(x) = k(x, x) - k_X(x)^\top K_{XX}^{-1}k_X(x)$$

Noisy observations/with nugget - Regression

In practice, we don't usually observe $f(x)$ directly. If we observe

$$y_i = f(x_i) + N(0, \sigma^2)$$

Noisy observations/with nugget - Regression

In practice, we don't usually observe $f(x)$ directly. If we observe

$$y_i = f(x_i) + N(0, \sigma^2)$$

then $y_1, \dots, y_n, f(x) \sim N_{n+1}(0, \Sigma)$

where $\Sigma = \begin{pmatrix} K_{xx} + \sigma^2 I & | & k(x_1, x) \\ k(x_2, x) & | & k(x_2, x) \\ \vdots & | & k(x_n, x) \\ \hline k(x, x_1) & k(x, x_2) & \dots & k(x, x_n) & | & k(x, x) \end{pmatrix}$

Noisy observations/with nugget - Regression

In practice, we don't usually observe $f(x)$ directly. If we observe

$$y_i = f(x_i) + N(0, \sigma^2)$$

then $y_1, \dots, y_n, f(x) \sim N_{n+1}(0, \Sigma)$

where $\Sigma = \begin{pmatrix} K_{XX} + \sigma^2 I & | & k(x_1, x) \\ k(x_2, x) & | & k(x_2, x) \\ \vdots & | & \vdots \\ k(x_n, x) & | & k(x_n, x) \end{pmatrix}$

$$\Sigma = \begin{pmatrix} K_{XX} + \sigma^2 I & | & k(x_1, x) \\ k(x_2, x) & | & k(x_2, x) \\ \vdots & | & \vdots \\ k(x_n, x) & | & k(x_n, x) \end{pmatrix}$$

Then

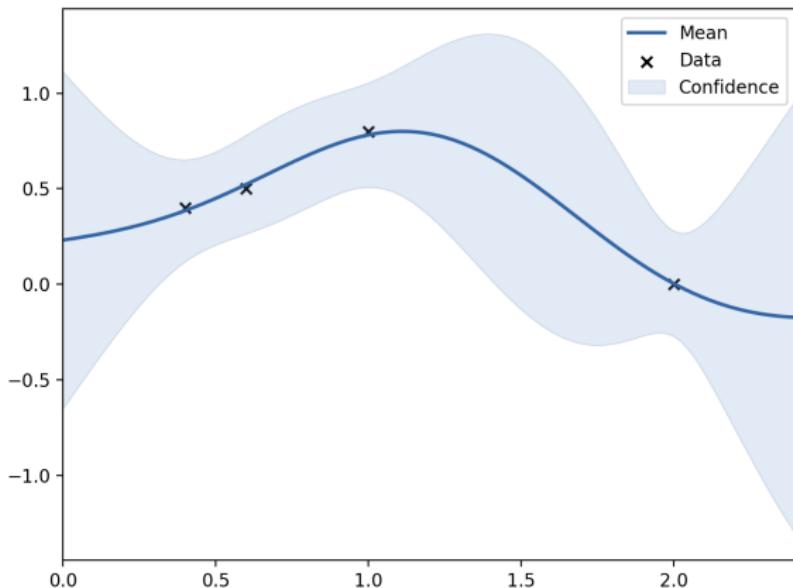
$$f(x) \mid y_1, \dots, y_n \sim N(\bar{m}(x), \bar{k}(x))$$

where

$$\bar{m}(x) = m(x) + k_X(x)^\top (K_{XX} + \sigma^2 I)^{-1} (\mathbf{y} - m_X)$$

$$\bar{k}(x) = k(x, x) - k_X(x)^\top (K_{XX} + \sigma^2 I)^{-1} k_X(x)$$

Nugget standard deviation $\sigma = 0.1$

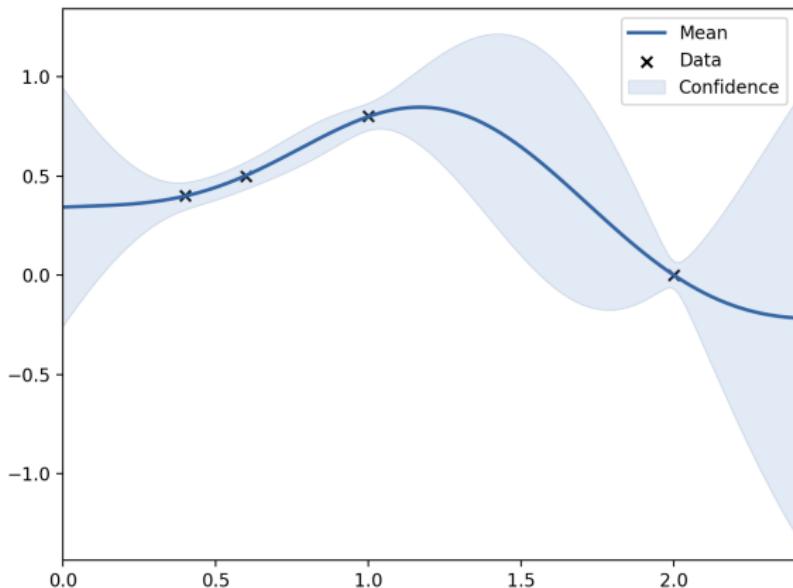


Solid line $\bar{m}(x) = m(x) + k_X(x)^\top K_{XX}^{-1}(\mathbf{y} - m_X)$

Shaded region $\bar{m}(x) \pm 1.96\sqrt{\bar{k}(x)}$

$$\bar{k}(x) = k(x, x) - k_X(x)^\top (K_{XX}^{-1} + \sigma^2 I) k_X(x)$$

Nugget standard deviation $\sigma = 0.025$



Solid line $\bar{m}(x) = m(x) + k_X(x)^\top K_{XX}^{-1}(\mathbf{y} - m_X)$

Shaded region $\bar{m}(x) \pm 1.96\sqrt{\bar{k}(x)}$

$$\bar{k}(x) = k(x, x) - k_X(x)^\top (K_{XX}^{-1} + \sigma^2 I) k_X(x)$$

Why use GPs? Answer 1

The GP class of models is closed under various operations.

Why use GPs? Answer 1

The GP class of models is closed under various operations.

- Closed under Bayesian conditioning, i.e., if we observe

$$\mathbf{D} = (f(x_1), \dots, f(x_n))$$

then

$$f | D \sim GP$$

but with updated mean and covariance functions.

Why use GPs? Answer 1

The GP class of models is closed under various operations.

- Closed under Bayesian conditioning, i.e., if we observe

$$\mathbf{D} = (f(x_1), \dots, f(x_n))$$

then

$$f | D \sim GP$$

but with updated mean and covariance functions.

- Closed under addition

$$f_1(\cdot), f_2(\cdot) \sim GP \quad \text{then} \quad (f_1 + f_2)(\cdot) \sim GP$$

Why use GPs? Answer 1

The GP class of models is closed under various operations.

- Closed under Bayesian conditioning, i.e., if we observe

$$\mathbf{D} = (f(x_1), \dots, f(x_n))$$

then

$$f | D \sim GP$$

but with updated mean and covariance functions.

- Closed under addition

$$f_1(\cdot), f_2(\cdot) \sim GP \quad \text{then} \quad (f_1 + f_2)(\cdot) \sim GP$$

- Closed under any linear operator. If $f \sim GP(m(\cdot), k(\cdot, \cdot))$, then if \mathcal{L} is a linear operator

$$\mathcal{L} \circ f \sim GP(\mathcal{L} \circ m, \mathcal{L}^2 \circ k)$$

e.g. $\frac{df}{dx}, \int f(x)dx, Af$ are all GPs

Motivation for non-parametric models

Functions live in function spaces (vector spaces with inner products). There are lots of different function spaces: the GP kernel k implicitly determines a particular (RKHS) function/hypothesis space

- Generally, we don't think too hard about this space, we just choose a kernel and attempt to validate it empirically.

³and can be dense in some sets of continuous bounded functions

Motivation for non-parametric models

Functions live in function spaces (vector spaces with inner products). There are lots of different function spaces: the GP kernel k implicitly determines a particular (RKHS) function/hypothesis space

- Generally, we don't think too hard about this space, we just choose a kernel and attempt to validate it empirically.

Although reality may not lie in the RKHS defined by k , this space is much richer than any parametric regression model³,

- thus is more likely to contain a function close to the truth than any finite dimensional class of functions (eg linear models)

³and can be dense in some sets of continuous bounded functions

Hyperparameter estimation

If we know the covariance function, GPs work great!

Hyperparameter estimation

If we know the covariance function, GPs work great!

- Unfortunately, we usually don't! We need to estimate

Hyperparameter estimation

If we know the covariance function, GPs work great!

- Unfortunately, we usually don't! We need to estimate

A typical GP workflow:

- ① Pick a covariance function k from a small set, based on differentiability considerations.

Hyperparameter estimation

If we know the covariance function, GPs work great!

- Unfortunately, we usually don't! We need to estimate

A typical GP workflow:

- ➊ Pick a covariance function k from a small set, based on differentiability considerations.
- ➋ Covariance functions usually contain hyper-parameters ψ .

$$\text{RBF kernel } k(x, x') = \sigma^2 \exp\left(-\frac{1}{2} \frac{(x - x')^2}{\lambda^2}\right)$$

where $\psi = (\sigma^2, \lambda)$.

Hyperparameter estimation

If we know the covariance function, GPs work great!

- Unfortunately, we usually don't! We need to estimate

A typical GP workflow:

- ➊ Pick a covariance function k from a small set, based on differentiability considerations.
- ➋ Covariance functions usually contain hyper-parameters ψ .

$$\text{RBF kernel } k(x, x') = \sigma^2 \exp\left(-\frac{1}{2} \frac{(x - x')^2}{\lambda^2}\right)$$

where $\psi = (\sigma^2, \lambda)$.

We estimate ψ using your favourite statistical procedure.

- ▶ E.g. maximum likelihood:

$$\hat{\psi} = \arg \max_{\psi} \log \pi(\mathbf{y}|\mathbf{X}, \psi) = \arg \max_{\psi} \log N(y|m(X), K_{XX})$$

Hyperparameter estimation

If we know the covariance function, GPs work great!

- Unfortunately, we usually don't! We need to estimate

A typical GP workflow:

- ➊ Pick a covariance function k from a small set, based on differentiability considerations.
- ➋ Covariance functions usually contain hyper-parameters ψ .

$$\text{RBF kernel } k(x, x') = \sigma^2 \exp\left(-\frac{1}{2} \frac{(x - x')^2}{\lambda^2}\right)$$

where $\psi = (\sigma^2, \lambda)$.

We estimate ψ using your favourite statistical procedure.

- ▶ E.g. maximum likelihood:

$$\hat{\psi} = \arg \max_{\psi} \log \pi(\mathbf{y}|\mathbf{X}, \psi) = \arg \max_{\psi} \log N(y|m(X), K_{XX})$$

- ➌ Possibly try a few k (plus a few combinations), and use empirical evaluation to make a good choice.

Code

Most GP software essentially follows these steps.

Code

Most GP software essentially follows these steps.

- ① Pick k

```
gpt.kernels.ScaleKernel(gpt.kernels.RBFKernel())
```

Code

Most GP software essentially follows these steps.

- ① Pick k

```
gpt.kernels.ScaleKernel(gpt.kernels.RBFKernel())
```

- ② Pick m

```
gpt.means.ZeroMean()
```

Code

Most GP software essentially follows these steps.

- ① Pick k

```
gpt.kernels.ScaleKernel(gpt.kernels.RBFKernel())
```

- ② Pick m

```
gpt.means.ZeroMean()
```

- ③ Define the model

```
likelihood = gpt.likelihoods.GaussianLikelihood()  
model = ExactGPModel(X, y, likelihood)
```

Code

Most GP software essentially follows these steps.

- ① Pick k

```
gpt.kernels.ScaleKernel(gpt.kernels.RBFKernel())
```

- ② Pick m

```
gpt.means.ZeroMean()
```

- ③ Define the model

```
likelihood = gpt.likelihoods.GaussianLikelihood()
```

```
model = ExactGPModel(X, y, likelihood)
```

- ④ Optimize the hyper parameters;

```
optimizer = torch.optim.Adam(model.parameters(), lr=0.1)
mll = gpt.mlls.ExactMarginalLogLikelihood(likelihood, model)
for i in range(1000):
    output = model(X)
    loss = -mll(output, y)
    loss.backward()
    optimizer.step()
```

Software

This is the syntax for GPyTorch

Software

This is the syntax for GPyTorch

- Uses Torch to exploit auto-diff capabilities, and easy deployment on a GPU
- Powerful and easily adaptable
- Slightly harder to learn than some other packages

Software

This is the syntax for GPyTorch

- Uses Torch to exploit auto-diff capabilities, and easy deployment on a GPU
- Powerful and easily adaptable
- Slightly harder to learn than some other packages

There are many other Python (and R, Julia., Matlab,...) GP packages

- GPy - easy to use, no longer supported, hard to modify
- GPerks - easy to use, but limited functionality and hard to modify
- GPflow - uses tensorflow and is similar to GPyTorch
- ...

Benefits of GPs: Uncertainty estimates

GP predictions consist of two parts:

- point estimate $m(x) = \mathbb{E}f(x)$
- uncertainty about the estimate $k(x, x) = \text{Var}f(x)$

Quantification of prediction uncertainty (cf NNs) is one of the main advantages of GPs.

Benefits of GPs: Uncertainty estimates

GP predictions consist of two parts:

- point estimate $m(x) = \mathbb{E}f(x)$
- uncertainty about the estimate $k(x, x) = \text{Var}f(x)$

Quantification of prediction uncertainty (cf NNs) is one of the main advantages of GPs.

It is important to validate both aspects.

Split the data into training and test sets.

Benefits of GPs: Uncertainty estimates

GP predictions consist of two parts:

- point estimate $m(x) = \mathbb{E}f(x)$
- uncertainty about the estimate $k(x, x) = \text{Var}f(x)$

Quantification of prediction uncertainty (cf NNs) is one of the main advantages of GPs.

It is important to validate both aspects.

Split the data into training and test sets.

Then evaluate:

- Something that assess the point estimate on the test data, e.g., root mean square error (RMSE). Sometimes useful to convert this to an R^2 skill-score.

Benefits of GPs: Uncertainty estimates

GP predictions consist of two parts:

- point estimate $m(x) = \mathbb{E}f(x)$
- uncertainty about the estimate $k(x, x) = \text{Var}f(x)$

Quantification of prediction uncertainty (cf NNs) is one of the main advantages of GPs.

It is important to validate both aspects.

Split the data into training and test sets.

Then evaluate:

- Something that assess the point estimate on the test data, e.g., root mean square error (RMSE). Sometimes useful to convert this to an R^2 skill-score.
- Something that assesses the uncertainty quantification:
 - ▶ Coverage: e.g., what % of predictions lie in the GP 90% credible interval?

Benefits of GPs: Uncertainty estimates

GP predictions consist of two parts:

- point estimate $m(x) = \mathbb{E}f(x)$
- uncertainty about the estimate $k(x, x) = \text{Var}f(x)$

Quantification of prediction uncertainty (cf NNs) is one of the main advantages of GPs.

It is important to validate both aspects.

Split the data into training and test sets.

Then evaluate:

- Something that assess the point estimate on the test data, e.g., root mean square error (RMSE). Sometimes useful to convert this to an R^2 skill-score.
- Something that assesses the uncertainty quantification:
 - ▶ Coverage: e.g., what % of predictions lie in the GP 90% credible interval?
- Predictive log-likelihood (or another proper score) - assesses both aspects (point and uncertainty) but hard to interpret - best for comparing different covariance functions.

Problems with hyper-parameter optimization

The likelihood surface that is maximized in hyper-parameter estimation is often flat and multi-modal,

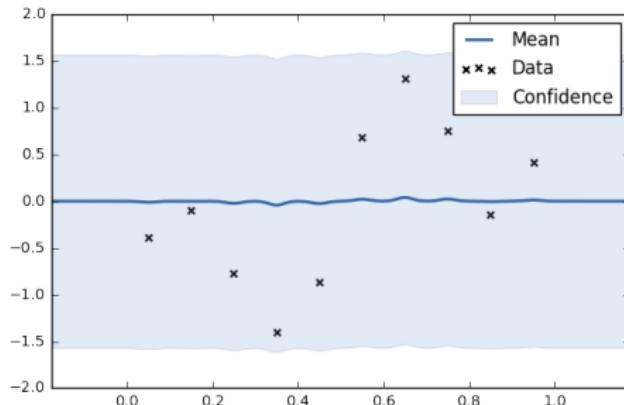
- optimizers can sometimes fail to converge, or get stuck in local-maxima.

Problems with hyper-parameter optimization

The likelihood surface that is maximized in hyper-parameter estimation is often flat and multi-modal,

- optimizers can sometimes fail to converge, or get stuck in local-maxima.

In practice, it is not uncommon to optimize hyper parameters and find solutions such as

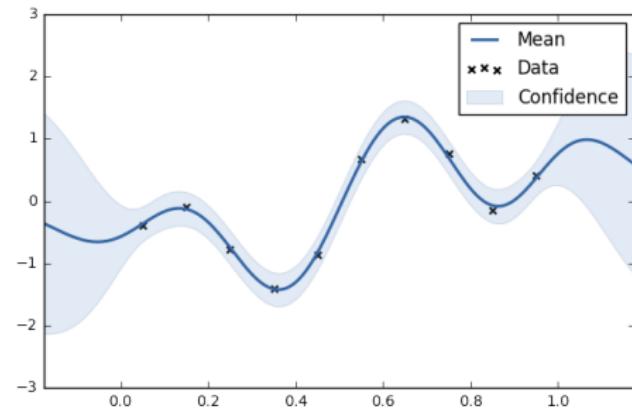
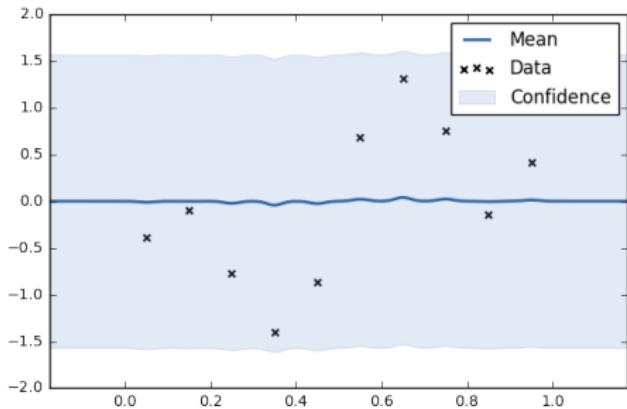


Problems with hyper-parameter optimization

The likelihood surface that is maximized in hyper-parameter estimation is often flat and multi-modal,

- optimizers can sometimes fail to converge, or get stuck in local-maxima.

In practice, it is not uncommon to optimize hyper parameters and find solutions such as



Work around these problems by running the optimizer multiple times from random start points, using prior distributions, constraining or fixing hyper-parameters, or adding white noise.

Computational cost

The computational cost of GP training is $O(n^3)$ with $O(n^2)$ memory requirements

Computational cost

The computational cost of GP training is $O(n^3)$ with $O(n^2)$ memory requirements

There are many ways to reduce these costs. E.g. basis expansions:

Computational cost

The computational cost of GP training is $O(n^3)$ with $O(n^2)$ memory requirements

There are many ways to reduce these costs. E.g. basis expansions:
Suppose

$$k(x, x') = \sum_{i=1}^m \phi_i(x)\phi_i(x') = \phi(x)^\top \phi(x')$$

Computational cost

The computational cost of GP training is $O(n^3)$ with $O(n^2)$ memory requirements

There are many ways to reduce these costs. E.g. basis expansions:
Suppose

$$k(x, x') = \sum_{i=1}^m \phi_i(x)\phi_i(x') = \phi(x)^\top \phi(x')$$

Then GP regression is equivalent to linear regression with covariates $\phi(x)$

- Dual form for regression coefficients costs $O(n^3)$,
but primal solution only costs $O(m^3)$

Computational cost

The computational cost of GP training is $O(n^3)$ with $O(n^2)$ memory requirements

There are many ways to reduce these costs. E.g. basis expansions:
Suppose

$$k(x, x') = \sum_{i=1}^m \phi_i(x)\phi_i(x') = \phi(x)^\top \phi(x')$$

Then GP regression is equivalent to linear regression with covariates $\phi(x)$

- Dual form for regression coefficients costs $O(n^3)$,
but primal solution only costs $O(m^3)$

It is possible now to do GP regression with $\sim 10^{6-7}$ data points.

- But GPs are primarily used in the low-data regime ($n \leq 10^4$), with neural networks a more popular choice if big datasets are available.

Curse of dimensionality

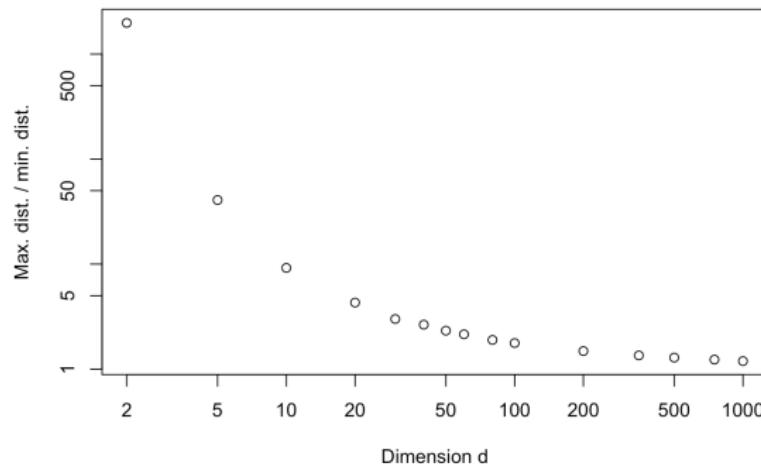
As $d = \dim(x)$ increases, regression gets much harder

- A d -dimensional cube $[0, 1]^d$ has 2^d corners: $2^{30} \approx 10^9$.

Curse of dimensionality

As $d = \dim(x)$ increases, regression gets much harder

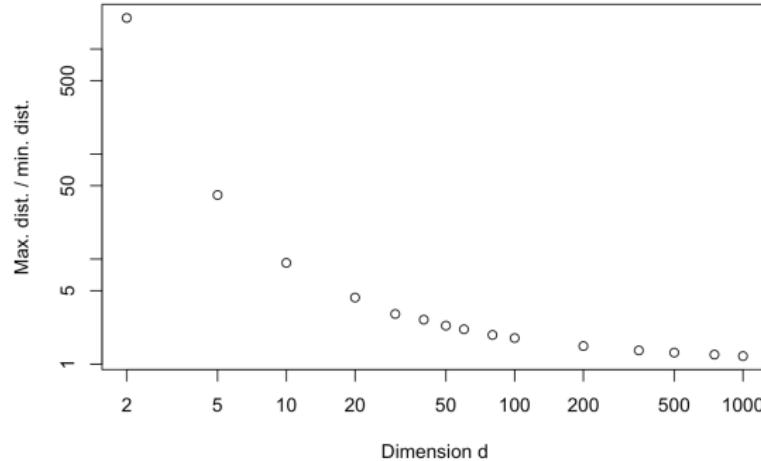
- A d -dimensional cube $[0, 1]^d$ has 2^d corners: $2^{30} \approx 10^9$.
- GPs work by computing distances between training points - for high d all points are approximately equidistant. Eg, 1000 points randomly chosen in the unit cube;



Curse of dimensionality

As $d = \dim(x)$ increases, regression gets much harder

- A d -dimensional cube $[0, 1]^d$ has 2^d corners: $2^{30} \approx 10^9$.
- GPs work by computing distances between training points - for high d all points are approximately equidistant. Eg, 1000 points randomly chosen in the unit cube;



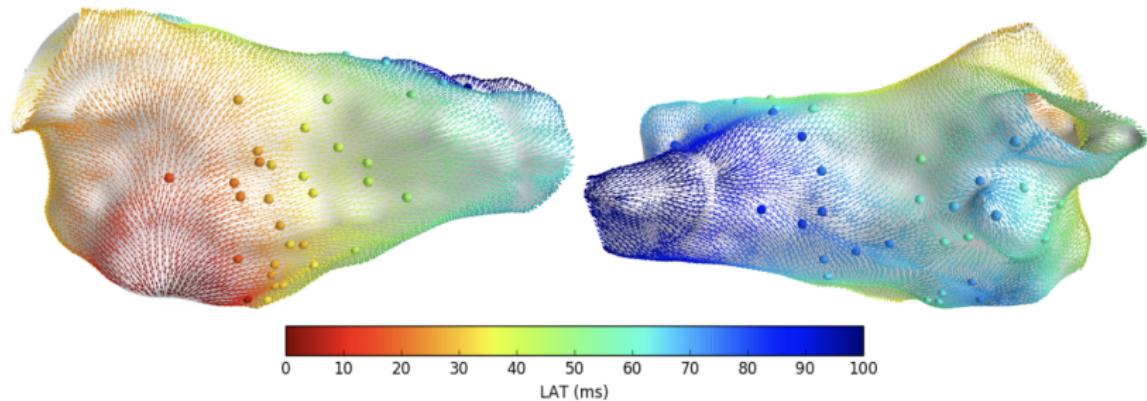
Can use dimension reductions to try to find informative low-dimensional manifolds: PCA, sensitivity analysis, random projection...

Example of GPs in Cardiac Digital Twins: Interpolating LATs

Coveney *et al.* 2019, 2020

We want to estimate activation times at all locations on the atria (the *LAT map*)

- Typically, only able to measure LAT a small number (~ 10) of locations on the atrium.

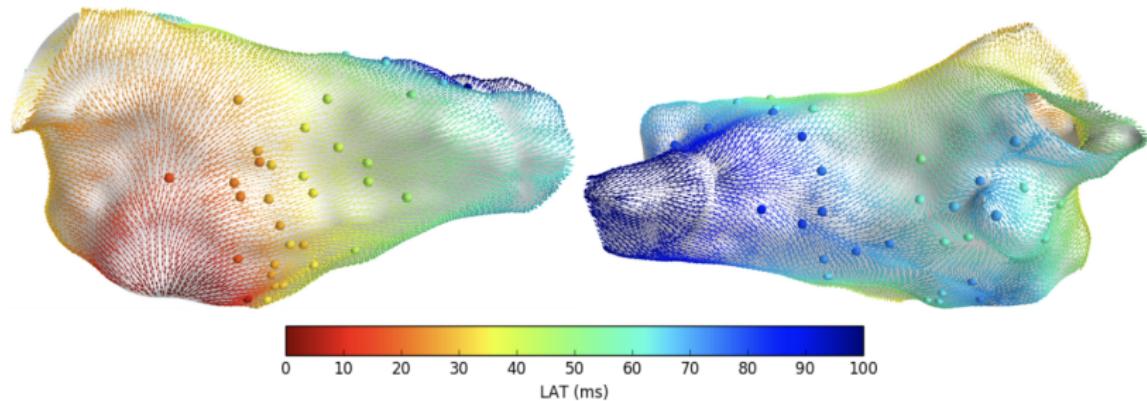


Example of GPs in Cardiac Digital Twins: Interpolating LATs

Coveney *et al.* 2019, 2020

We want to estimate activation times at all locations on the atria (the *LAT map*)

- Typically, only able to measure LAT a small number (~ 10) of locations on the atrium.



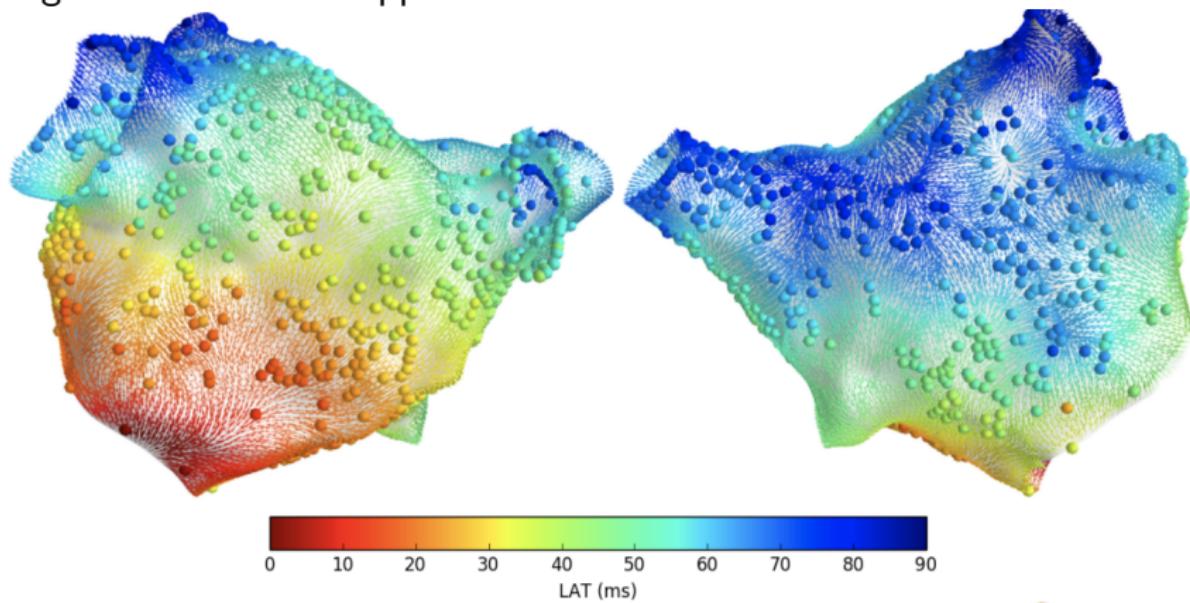
How can we interpolate to other locations and estimate conduction velocities?

$$LAT_{obs}(x) = LAT_{true}(x) + \epsilon$$

Results

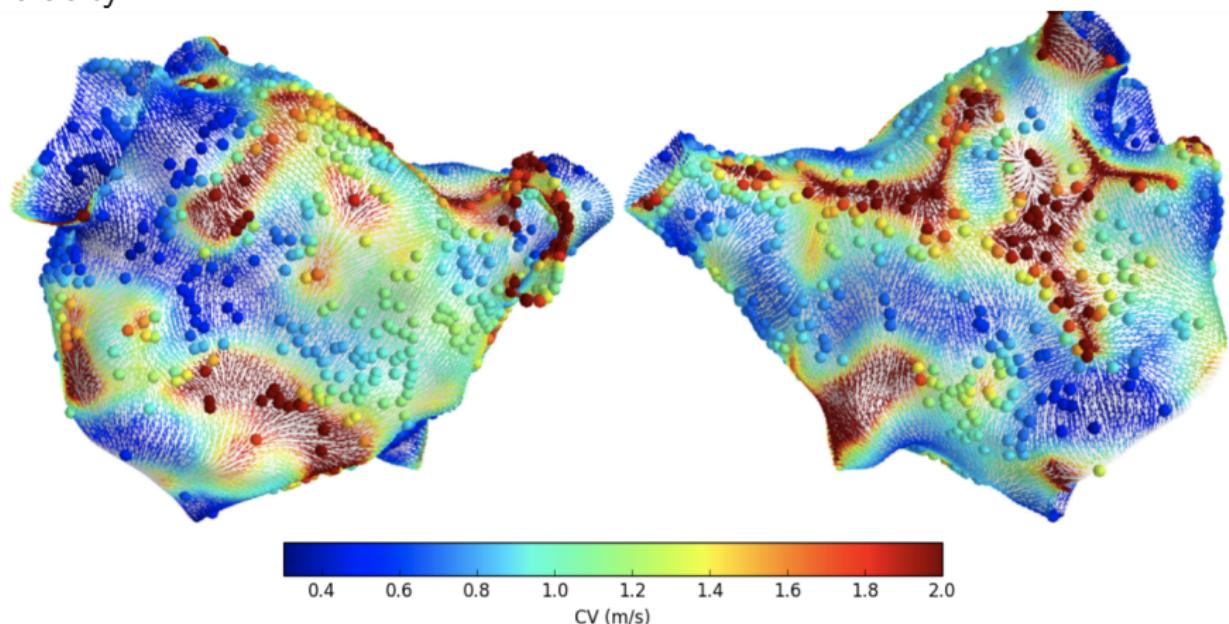
Coveney *et al.* 2019, 2020

Using a basis function approach we can model on the atrial manifold:



Results

Because GPs are differentiable, the approach easily gives conduction velocity:



Conclusions

- GPs are now ubiquitous in statistics/ML, and a key part of the cardiac digital twin pipeline.
- Popularity stems from
 - ▶ Naturalness of the framework
 - ▶ Mathematical tractability
 - ▶ Empirical success

Conclusions

- GPs are now ubiquitous in statistics/ML, and a key part of the cardiac digital twin pipeline.
- Popularity stems from
 - ▶ Naturalness of the framework
 - ▶ Mathematical tractability
 - ▶ Empirical success

Thank you for listening!

References

- Rasmussen and Williams. *Gaussian processes for machine learning.* MIT press, 2006.
- Stein. *Interpolation of Spatial Data: Some Theory for Kriging.* Springer, 1999
- Kanagawa, Hennig, Sejdinovic, and Sriperumbudur. *Gaussian processes and kernel methods: A review on connections and equivalences..* arXiv:1807.02582 2018.
- Many tutorial talks available at gpss.cc

Why use GPs? Answer 3: Naturalness of GP framework

Why use **Gaussian** processes as non-parametric models?

Why use GPs? Answer 3: Naturalness of GP framework

Why use **Gaussian** processes as non-parametric models?

If we only knew the expectation and variance of some random variables, X and Y , then how should we best do statistics?

Why use GPs? Answer 3: Naturalness of GP framework

Why use **Gaussian** processes as non-parametric models?

If we only knew the expectation and variance of some random variables, X and Y , then how should we best do statistics?

It has been shown, using coherency arguments, or geometric arguments, or..., that the best second-order inference we can do is to update our beliefs about X given Y using

$$\mathbb{E}(X|Y) = \mathbb{E}(X) + \text{Cov}(X, Y)\text{Var}(Y)^{-1}(Y - \mathbb{E}(Y))$$

i.e., exactly the Gaussian process update for the posterior mean.

So GPs are in some sense second-order optimal.

Kriging

Kriging

Suppose $Y(x)$ is a (second order stationary) stochastic process with

$$\begin{aligned}\mathbb{E} Y(x) &= \mu \quad \forall x \\ \text{Cov}(Y(x), Y(x')) &= k(x - x') \quad \forall x, x'\end{aligned}$$

NB we're not assuming Y has a Gaussian distribution.

Kriging

Suppose $Y(x)$ is a (second order stationary) stochastic process with

$$\begin{aligned}\mathbb{E} Y(x) &= \mu \quad \forall x \\ \text{Cov}(Y(x), Y(x')) &= k(x - x') \quad \forall x, x'\end{aligned}$$

NB we're not assuming Y has a Gaussian distribution.

If someone tells you $\mathbf{y} = (Y(x_1), \dots, Y(x_n))^\top$, how would you predict $Y(x)$?

Kriging

Suppose $Y(x)$ is a (second order stationary) stochastic process with

$$\begin{aligned}\mathbb{E} Y(x) &= \mu \quad \forall x \\ \text{Cov}(Y(x), Y(x')) &= k(x - x') \quad \forall x, x'\end{aligned}$$

NB we're not assuming Y has a Gaussian distribution.

If someone tells you $\mathbf{y} = (Y(x_1), \dots, Y(x_n))^\top$, how would you predict $Y(x)$?

One option is to find the best linear unbiased predictor (BLUP) of $Y(x)$.

Best Linear Unbiased Predictors (BLUP)

Consider the linear estimator

$$\hat{Y}(x) = c + \sum w_i Y(x_i) = c + \mathbf{w}^\top \mathbf{y}$$

Best Linear Unbiased Predictors (BLUP)

Consider the linear estimator

$$\hat{Y}(x) = c + \sum w_i Y(x_i) = c + \mathbf{w}^\top \mathbf{y}$$

If we require $\hat{Y}(x)$ to be unbiased,

$$\begin{aligned}\mu &= \mathbb{E}\hat{Y}(x) \\ &= \mathbb{E}(c + \mathbf{w}^\top \mathbf{y}) \\ &= c + \mathbf{w}^\top \boldsymbol{\mu}\end{aligned}$$

where $\boldsymbol{\mu} = (\mu, \dots, \mu)^\top$.

Best Linear Unbiased Predictors (BLUP)

Consider the linear estimator

$$\hat{Y}(x) = c + \sum w_i Y(x_i) = c + \mathbf{w}^\top \mathbf{y}$$

If we require $\hat{Y}(x)$ to be unbiased,

$$\begin{aligned}\mu &= \mathbb{E}\hat{Y}(x) \\ &= \mathbb{E}(c + \mathbf{w}^\top \mathbf{y}) \\ &= c + \mathbf{w}^\top \boldsymbol{\mu}\end{aligned}$$

where $\boldsymbol{\mu} = (\mu, \dots, \mu)^\top$.

Thus $c = \mu - \mathbf{w}^\top \boldsymbol{\mu}$ and we must have

$$\hat{Y}(x) = \mu + \mathbf{w}^\top (\mathbf{y} - \boldsymbol{\mu})$$

Best Linear Unbiased Predictors (BLUP) - II

The **best** linear unbiased predictor minimises the mean square error

$$\begin{aligned}MSE(\hat{Y}(x)) &= \mathbb{E}((\hat{Y}(x) - Y(x))^2) \\&= \mathbb{E}\left((\mathbf{w}^\top(\mathbf{y} - \boldsymbol{\mu}) + (\boldsymbol{\mu} - Y(x))^2)\right) \\&= \mathbf{w}^\top \text{Var}(\mathbf{y}) \mathbf{w} + \text{Var}(Y(x)) - 2\mathbf{w}^\top \text{Cov}(\mathbf{y}, Y(x)) \\&= \mathbf{w}^\top K_{XX} \mathbf{w} + k(0) - 2\mathbf{w}^\top \mathbf{k}_X(x)\end{aligned}$$

Best Linear Unbiased Predictors (BLUP) - II

The **best** linear unbiased predictor minimises the mean square error

$$\begin{aligned}MSE(\hat{Y}(x)) &= \mathbb{E}((\hat{Y}(x) - Y(x))^2) \\&= \mathbb{E}\left((\mathbf{w}^\top(\mathbf{y} - \boldsymbol{\mu}) + (\boldsymbol{\mu} - Y(x))^2)\right) \\&= \mathbf{w}^\top \text{Var}(\mathbf{y}) \mathbf{w} + \text{Var}(Y(x)) - 2\mathbf{w}^\top \text{Cov}(\mathbf{y}, Y(x)) \\&= \mathbf{w}^\top K_{XX} \mathbf{w} + k(0) - 2\mathbf{w}^\top \mathbf{k}_X(x)\end{aligned}$$

If we differentiate wrt w and set the gradient equal to zero, we find

$$0 = 2K_{XX} \mathbf{w} - 2\mathbf{k}_X(x)$$

Best Linear Unbiased Predictors (BLUP) - II

The **best** linear unbiased predictor minimises the mean square error

$$\begin{aligned}MSE(\hat{Y}(x)) &= \mathbb{E}((\hat{Y}(x) - Y(x))^2) \\&= \mathbb{E}\left((\mathbf{w}^\top(\mathbf{y} - \boldsymbol{\mu}) + (\boldsymbol{\mu} - Y(x))^2)\right) \\&= \mathbf{w}^\top \text{Var}(\mathbf{y}) \mathbf{w} + \text{Var}(Y(x)) - 2\mathbf{w}^\top \text{Cov}(\mathbf{y}, Y(x)) \\&= \mathbf{w}^\top K_{XX} \mathbf{w} + k(0) - 2\mathbf{w}^\top \mathbf{k}_X(x)\end{aligned}$$

If we differentiate wrt w and set the gradient equal to zero, we find

$$0 = 2K_{XX}\mathbf{w} - 2\mathbf{k}_X(x)$$

and thus

$$\hat{Y}(x) = \boldsymbol{\mu} + \mathbf{k}_X(x)^\top K_{XX}^{-1}(\mathbf{y} - \boldsymbol{\mu})$$

as before.

So the Gaussian process posterior mean is optimal (i.e. is the BLUP) even if we don't assume Gaussianity.

Why use GPs? Answer 2: non-parametric/kernel regression

We can also view GPs as a non-parametric extension to linear regression.

- k determines the space of functions that sample paths live in.

Why use GPs? Answer 2: non-parametric/kernel regression

We can also view GPs as a non-parametric extension to linear regression.

- k determines the space of functions that sample paths live in.

Suppose we're given data $\{(x_i, y_i)_{i=1}^n\}$ with $x_i \in \mathbb{R}^p, y_i \in \mathbb{R}$

Why use GPs? Answer 2: non-parametric/kernel regression

We can also view GPs as a non-parametric extension to linear regression.

- k determines the space of functions that sample paths live in.

Suppose we're given data $\{(x_i, y_i)_{i=1}^n\}$ with $x_i \in \mathbb{R}^p, y_i \in \mathbb{R}$

$$\hat{\beta} = \arg \min_{\beta} \|y - X\beta\|_2^2 + \sigma^2 \|\beta\|_2^2 \quad \text{regularised least squares}$$

where $X = \begin{pmatrix} x_1^\top \\ x_2^\top \\ \vdots \\ x_n^\top \end{pmatrix}$

Why use GPs? Answer 2: non-parametric/kernel regression

We can also view GPs as a non-parametric extension to linear regression.

- k determines the space of functions that sample paths live in.

Suppose we're given data $\{(x_i, y_i)_{i=1}^n\}$ with $x_i \in \mathbb{R}^p, y_i \in \mathbb{R}$

$$\begin{aligned}\hat{\beta} &= \arg \min_{\beta} \|y - X\beta\|_2^2 + \sigma^2 \|\beta\|_2^2 \quad \text{regularised least squares} \\ &= (X^\top X + \sigma^2 I)^{-1} X^\top y \quad \text{usual ridge regression estimator}\end{aligned}$$

where $X = \begin{pmatrix} x_1^\top \\ x_2^\top \\ \vdots \\ x_n^\top \end{pmatrix}$

Why use GPs? Answer 2: non-parametric/kernel regression

We can also view GPs as a non-parametric extension to linear regression.

- k determines the space of functions that sample paths live in.

Suppose we're given data $\{(x_i, y_i)_{i=1}^n\}$ with $x_i \in \mathbb{R}^p, y_i \in \mathbb{R}$

$$\begin{aligned}\hat{\beta} &= \arg \min_{\beta} \|y - X\beta\|_2^2 + \sigma^2 \|\beta\|_2^2 \quad \text{regularised least squares} \\ &= (X^\top X + \sigma^2 I)^{-1} X^\top y \quad \text{usual ridge regression estimator} \\ &= X^\top (X X^\top + \sigma^2 I)^{-1} y \quad \text{the dual form}\end{aligned}$$

where $X = \begin{pmatrix} x_1^\top \\ x_2^\top \\ \vdots \\ x_n^\top \end{pmatrix}$

Why use GPs? Answer 2: non-parametric/kernel regression

We can also view GPs as a non-parametric extension to linear regression.

- k determines the space of functions that sample paths live in.

Suppose we're given data $\{(x_i, y_i)_{i=1}^n\}$ with $x_i \in \mathbb{R}^p, y_i \in \mathbb{R}$

$$\hat{\beta} = \arg \min_{\beta} \|y - X\beta\|_2^2 + \sigma^2 \|\beta\|_2^2 \quad \text{regularised least squares}$$

$$= (X^\top X + \sigma^2 I)^{-1} X^\top y \quad \text{usual ridge regression estimator}$$

$$= X^\top (X X^\top + \sigma^2 I)^{-1} y \quad \text{the dual form}$$

$$\text{as } (X^\top X + \sigma^2 I) X^\top = X^\top (X X^\top + \sigma^2 I)$$

$$\text{so } X^\top (X X^\top + \sigma^2 I)^{-1} = (X^\top X + \sigma^2 I)^{-1} X^\top$$

$$\text{where } X = \begin{pmatrix} x_1^\top \\ x_2^\top \\ \vdots \\ x_n^\top \end{pmatrix}$$

At first the dual form

$$\hat{\beta} = X^\top (X X^\top + \sigma^2 I)^{-1} y$$

looks harder to compute than the usual

$$\hat{\beta} = (X^\top X + \sigma^2 I)^{-1} X^\top y$$

- $X^\top X$ is $p \times p$ p = number of features/parameters
- XX^\top is $n \times n$ n is the number of data points

At first the dual form

$$\hat{\beta} = X^\top (X X^\top + \sigma^2 I)^{-1} y$$

looks harder to compute than the usual

$$\hat{\beta} = (X^\top X + \sigma^2 I)^{-1} X^\top y$$

- $X^\top X$ is $p \times p$ p = number of features/parameters
- $X X^\top$ is $n \times n$ n is the number of data points

But the dual form only uses inner products between vectors in \mathbb{R}^n

$$X X^\top = \begin{pmatrix} x_1^\top \\ \vdots \\ x_n^\top \end{pmatrix} (x_1 \dots x_n) = \begin{pmatrix} x_1^\top x_1 & \dots & x_1^\top x_n \\ \vdots & & \vdots \\ x_n^\top x_1 & \dots & x_n^\top x_n \end{pmatrix}$$

$= K_{XX}$ if $k(x, x') = x^\top x'$

— This is useful!

Prediction

The best prediction of y at a new location x' is

$$\begin{aligned}\hat{y}' &= x'^\top \hat{\beta} \\ &= x'^\top X^\top (X X^\top + \sigma^2 I)^{-1} y \\ &= k_X(x')^\top (K_{XX} + \sigma^2 I)^{-1} y\end{aligned}$$

where $k_X(x')^\top := (x'^\top x_1, \dots, x'^\top x_n)$ and $[K_{XX}]_{ij} := x_i^\top x_j$

Prediction

The best prediction of y at a new location x' is

$$\begin{aligned}\hat{y}' &= x'^\top \hat{\beta} \\ &= x'^\top X^\top (X X^\top + \sigma^2 I)^{-1} y \\ &= k_X(x')^\top (K_{XX} + \sigma^2 I)^{-1} y\end{aligned}$$

where $k_X(x')^\top := (x'^\top x_1, \dots, x'^\top x_n)$ and $[K_{XX}]_{ij} := x_i^\top x_j$
 K_{XX} and $k_X(x)$ are kernel matrices:

- every element is an inner product between 2 points: $k(x, x') = x^\top x'$

Prediction

The best prediction of y at a new location x' is

$$\begin{aligned}\hat{y}' &= x'^\top \hat{\beta} \\ &= x'^\top X^\top (X X^\top + \sigma^2 I)^{-1} y \\ &= k_X(x')^\top (K_{XX} + \sigma^2 I)^{-1} y\end{aligned}$$

where $k_X(x')^\top := (x'^\top x_1, \dots, x'^\top x_n)$ and $[K_{XX}]_{ij} := x_i^\top x_j$
 K_{XX} and $k_X(x)$ are kernel matrices:

- every element is an inner product between 2 points: $k(x, x') = x^\top x'$

Note this is the GP conditional mean we derived before with $m(x) = 0$.

$$m(x) = k_X(x)^\top (K_{XX} + \sigma^2 I)^{-1} y$$

- linear regression and GP regression are equivalent when $k(x, x') = x^\top x'$.

Including features I

We can replace x by a feature vector in linear regression, e.g.,
 $\phi(x) = (1 \ x \ x^2)$

It doesn't change the expressions other than the inner product if

$$k(x', x) = x'^\top x$$

is replaced by

$$k(x', x) = \phi(x')^\top \phi(x)$$

Including features I

We can replace x by a feature vector in linear regression, e.g.,
 $\phi(x) = (1 \ x \ x^2)$

It doesn't change the expressions other than the inner product if

$$k(x', x) = x'^\top x$$

is replaced by

$$k(x', x) = \phi(x')^\top \phi(x)$$

Note $k(x', x) = \phi(x')^\top \phi(x)$ is a semi-definite function for any choice of $\phi(x)$.

Including features II

For some sets of features, $\phi(x)$, computation of the inner product doesn't require us to evaluate the individual features.

Including features II

For some sets of features, $\phi(x)$, computation of the inner product doesn't require us to evaluate the individual features.

E.g., Consider $\mathcal{X} = \mathbb{R}^2$ and let

$$\phi : \mathbf{x} = (x_1, x_2) \mapsto (1, \sqrt{2}x_1, \sqrt{2}x_2, x_1^2, \sqrt{2}x_1x_2, x_2^2)^\top$$

i.e., linear regression using all the linear and quadratic terms, and first order interactions.

Including features II

For some sets of features, $\phi(x)$, computation of the inner product doesn't require us to evaluate the individual features.

E.g., Consider $\mathcal{X} = \mathbb{R}^2$ and let

$$\phi : \mathbf{x} = (x_1, x_2) \mapsto (1, \sqrt{2}x_1, \sqrt{2}x_2, x_1^2, \sqrt{2}x_1x_2, x_2^2)^\top$$

i.e., linear regression using all the linear and quadratic terms, and first order interactions.

Then

$$\begin{aligned} k(\mathbf{x}, \mathbf{z}) &= \phi(\mathbf{x})^\top \phi(\mathbf{z}) \\ &= (1, \sqrt{2}x_1, \sqrt{2}x_2, x_1^2, \sqrt{2}x_1x_2, x_2^2)(1, \sqrt{2}z_1, \sqrt{2}z_2, z_1^2, \sqrt{2}z_1z_2, z_2^2)^\top \\ &= (1 + (\mathbf{x}_1, \mathbf{x}_2)(\mathbf{z}_1, \mathbf{z}_2)^\top)^2 \\ &= (1 + \mathbf{x}^\top \mathbf{z})^2 \end{aligned}$$

Including features II

For some sets of features, $\phi(x)$, computation of the inner product doesn't require us to evaluate the individual features.

E.g., Consider $\mathcal{X} = \mathbb{R}^2$ and let

$$\phi : \mathbf{x} = (x_1, x_2) \mapsto (1, \sqrt{2}x_1, \sqrt{2}x_2, x_1^2, \sqrt{2}x_1x_2, x_2^2)^\top$$

i.e., linear regression using all the linear and quadratic terms, and first order interactions.

Then

$$\begin{aligned} k(\mathbf{x}, \mathbf{z}) &= \phi(\mathbf{x})^\top \phi(\mathbf{z}) \\ &= (1, \sqrt{2}x_1, \sqrt{2}x_2, x_1^2, \sqrt{2}x_1x_2, x_2^2)(1, \sqrt{2}z_1, \sqrt{2}z_2, z_1^2, \sqrt{2}z_1z_2, z_2^2)^\top \\ &= (1 + (\mathbf{x}_1, \mathbf{x}_2)(\mathbf{z}_1, \mathbf{z}_2)^\top)^2 \\ &= (1 + \mathbf{x}^\top \mathbf{z})^2 \end{aligned}$$

To evaluate $k(\mathbf{x}, \mathbf{z})$ we didn't need to explicitly compute the feature vector $\phi(\mathbf{x})$

Including features III

To evaluate $k(\mathbf{x}, \mathbf{z})$ we didn't need to explicitly compute the feature vectors $\phi(\mathbf{x}), \phi(\mathbf{z}) \in \mathbb{R}^6$

The same idea works with much larger feature vectors, sometimes even when $\phi(\mathbf{x}) \in \mathbb{R}^\infty$

Including features III

To evaluate $k(\mathbf{x}, \mathbf{z})$ we didn't need to explicitly compute the feature vectors $\phi(\mathbf{x}), \phi(\mathbf{z}) \in \mathbb{R}^6$

The same idea works with much larger feature vectors, sometimes even when $\phi(\mathbf{x}) \in \mathbb{R}^\infty$

Theorem: A function

$$k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$$

is positive semi-definite (and thus a valid covariance function) if and only if we can write

$$k(x, x') = \phi(x)^\top \phi(x')$$

for some (possibly infinite dimensional) feature vector $\phi(x)$.

Including features III

To evaluate $k(\mathbf{x}, \mathbf{z})$ we didn't need to explicitly compute the feature vectors $\phi(\mathbf{x}), \phi(\mathbf{z}) \in \mathbb{R}^6$

The same idea works with much larger feature vectors, sometimes even when $\phi(\mathbf{x}) \in \mathbb{R}^\infty$

Theorem: A function

$$k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$$

is positive semi-definite (and thus a valid covariance function) if and only if we can write

$$k(x, x') = \phi(x)^\top \phi(x')$$

for some (possibly infinite dimensional) feature vector $\phi(x)$.

So GP regression with k can be thought of as linear regression with $\phi(x)$.

Example: If $\mathcal{X} = \mathbb{R}$, set $\phi_c(x) = e^{-\frac{(x-c_0)^2}{2\lambda^2}}$ and

$$\phi^N(x) = \frac{1}{\sqrt{N}}(\phi_{c_0}(x), \dots, \phi_{c_N}(x))$$

with $c_0 = -\log N$, $c_N = \log N$, $c_{i+1} - c_i = 2\frac{\log N}{N}$.

Example: If $\mathcal{X} = \mathbb{R}$, set $\phi_c(x) = e^{-\frac{(x-c_0)^2}{2\lambda^2}}$ and

$$\phi^N(x) = \frac{1}{\sqrt{N}}(\phi_{c_0}(x), \dots, \phi_{c_N}(x))$$

with $c_0 = -\log N, c_N = \log N, c_{i+1} - c_i = 2\frac{\log N}{N}$. Then

$$\begin{aligned}\phi^N(x)^\top \phi^N(x') &= \frac{1}{N} \sum_{i=0}^N \phi_{c_i}(x) \phi_{c_i}(x') \\ &= \int_{-\log N}^{\log N} \phi_c(x) \phi(x') dc \rightarrow \exp\left(-\frac{(x-x')^2}{2\lambda^2}\right) \text{ as } N \rightarrow \infty\end{aligned}$$

So the RBF kernel arises when we do linear regression with an infinite feature vector containing Gaussian bumps

Example: If $\mathcal{X} = \mathbb{R}$, set $\phi_c(x) = e^{-\frac{(x-c_0)^2}{2\lambda^2}}$ and

$$\phi^N(x) = \frac{1}{\sqrt{N}}(\phi_{c_0}(x), \dots, \phi_{c_N}(x))$$

with $c_0 = -\log N$, $c_N = \log N$, $c_{i+1} - c_i = 2\frac{\log N}{N}$. Then

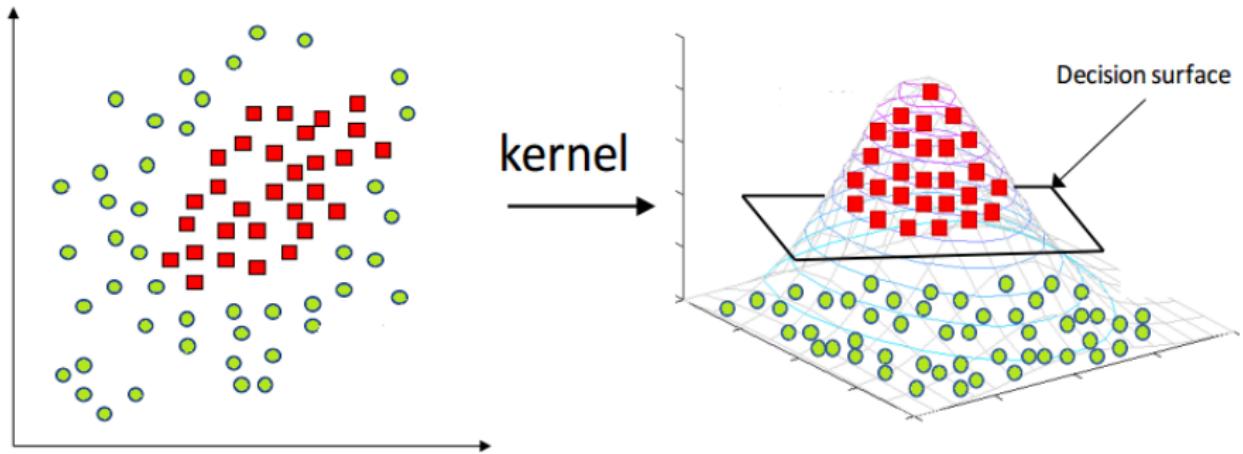
$$\begin{aligned}\phi^N(x)^\top \phi^N(x') &= \frac{1}{N} \sum_{i=0}^N \phi_{c_i}(x) \phi_{c_i}(x') \\ &= \int_{-\log N}^{\log N} \phi_c(x) \phi(x') dc \rightarrow \exp\left(-\frac{(x-x')^2}{2\lambda^2}\right) \text{ as } N \rightarrow \infty\end{aligned}$$

So the RBF kernel arises when we do linear regression with an infinite feature vector containing Gaussian bumps

We can use an infinite dimensional feature vector $\phi(x)$, and because linear regression can be done solely in terms of inner-products (inverting a $n \times n$ matrix in the dual form) we never need evaluate the feature vector, only the kernel.

Kernel trick:

lift x into feature space by replacing inner products $x^T x'$ by $k(x, x')$



Kernel regression (see Kanagawa et al. 2019)

Kernel regression and GP regression are closely related.

Kernel regression (see Kanagawa et al. 2019)

Kernel regression and GP regression are closely related.

Consider the space of functions

$$\mathcal{H}_k = \overline{\text{span}}\{k(\cdot, x) : x \in \mathcal{X}\}$$

ie functions of the form $\sum_{i=1}^n \alpha_i k(x, x_i)$ with inner product

$$\langle \sum a_i k(\cdot, x_i), \sum b_i k(\cdot, y_i) \rangle = \sum_{ij} a_i b_j k(x_i, y_j)$$

Kernel regression (see Kanagawa et al. 2019)

Kernel regression and GP regression are closely related.

Consider the space of functions

$$\mathcal{H}_k = \overline{\text{span}}\{k(\cdot, x) : x \in \mathcal{X}\}$$

ie functions of the form $\sum_{i=1}^n \alpha_i k(x, x_i)$ with inner product

$$\langle \sum a_i k(\cdot, x_i), \sum b_i k(\cdot, y_i) \rangle = \sum_{ij} a_i b_j k(x_i, y_j)$$

This is the reproducing kernel Hilbert space (RKHS) associated with k .

Kernel regression (see Kanagawa et al. 2019)

Kernel regression and GP regression are closely related.

Consider the space of functions

$$\mathcal{H}_k = \overline{\text{span}}\{k(\cdot, x) : x \in \mathcal{X}\}$$

ie functions of the form $\sum_{i=1}^n \alpha_i k(x, x_i)$ with inner product

$$\langle \sum a_i k(\cdot, x_i), \sum b_i k(\cdot, y_i) \rangle = \sum_{ij} a_i b_j k(x_i, y_j)$$

This is the reproducing kernel Hilbert space (RKHS) associated with k .

Kernel ridge regression chooses $f \in \mathcal{H}_k$ to minimise

$$L(f) = \sum_i (f(x_i) - y_i)^2 + \sigma^2 \|f\|_{\mathcal{H}_k}^2$$

Kernel regression (see Kanagawa et al. 2019)

Kernel regression and GP regression are closely related.

Consider the space of functions

$$\mathcal{H}_k = \overline{\text{span}}\{k(\cdot, x) : x \in \mathcal{X}\}$$

i.e. functions of the form $\sum_{i=1}^n \alpha_i k(x, x_i)$ with inner product

$$\langle \sum a_i k(\cdot, x_i), \sum b_i k(\cdot, y_i) \rangle = \sum_{ij} a_i b_j k(x_i, y_j)$$

This is the reproducing kernel Hilbert space (RKHS) associated with k .
Kernel ridge regression chooses $f \in \mathcal{H}_k$ to minimise

$$L(f) = \sum_i (f(x_i) - y_i)^2 + \sigma^2 \|f\|_{\mathcal{H}_k}^2$$

We can show that

$$\bar{m}(x) = \arg \min_{f \in \mathcal{H}_k} L(f)$$

where $\bar{m}(x)$ is the same as the posterior mean when we assume
 $y_i = f(x_i) + N(0, \sigma^2)$ and $f(\cdot) \sim GP(0, k(\cdot, \cdot))$

Kernel regression (see Kanagawa et al. 2019)

Kernel regression and GP regression are closely related.

Consider the space of functions

$$\mathcal{H}_k = \overline{\text{span}}\{k(\cdot, x) : x \in \mathcal{X}\}$$

i.e. functions of the form $\sum_{i=1}^n \alpha_i k(x, x_i)$ with inner product

$$\langle \sum a_i k(\cdot, x_i), \sum b_i k(\cdot, y_i) \rangle = \sum_{ij} a_i b_j k(x_i, y_j)$$

This is the reproducing kernel Hilbert space (RKHS) associated with k .

Kernel ridge regression chooses $f \in \mathcal{H}_k$ to minimise

$$L(f) = \sum_i (f(x_i) - y_i)^2 + \sigma^2 \|f\|_{\mathcal{H}_k}^2$$

We can show that

$$\bar{m}(x) = \arg \min_{f \in \mathcal{H}_k} L(f)$$

where $\bar{m}(x)$ is the same as the posterior mean when we assume $y_i = f(x_i) + N(0, \sigma^2)$ and $f(\cdot) \sim GP(0, k(\cdot, \cdot))$

Note: $\bar{m}(\cdot) \in \mathcal{H}_k$ but samples from a GP live in a slightly larger RKHS.